

Comparing Value-based and Policy-based Deep RL on a 2D Arcade Game and a 3D FPS

DQN vs PPO on Atari Pong and VizDoom Defend-the-Center / Health-Gathering

Michael Kurniawan Soegeng 吳忠賢—314540066

0.1 Abstract

This report studies which design choices most affect deep reinforcement-learning performance, and whether they transfer between a fully-observable 2D arcade game (Atari Pong) and a partially-observable 3D first-person shooter (VizDoom Defend-the-Center, Health-Gathering). Using Stable-Baselines3, identical 84×84 frame-stack image pipelines, and ≥ 3 seeds per curve, five threads are tested: (T1) value-based DQN vs policy-based PPO on Pong at an equal 7 M-step budget; (T2) DQN components on Pong —target network, ϵ -greedy schedule, replay buffer; (T3) frame stacking on the 3D task; (T4) Discrete vs MultiDiscrete action spaces; (T5) task complexity across two VizDoom scenarios. The clearest findings are (i) **DQN is markedly more sample-efficient than PPO on Pong at equal data** (+4.77 vs −6.58 / −6.07 at 7 M steps, where the second PPO number uses the literature SB3-zoo recipe —both PPO variants lose to DQN, so the inversion is not a hyperparameter artefact), (ii) **on Defend-the-Center the two are tied** at 2 M (DQN +8.85, PPO +9.37), suggesting the Pong gap is a sample-efficiency effect not a fundamental algorithm gap, (iii) **the lagged target network is the single most important DQN stability component** (≈ 9.5 -point gap at the 2 M ablation budget; replay-buffer size is only ≈ 0.9 pt), and (iv) **frame stacking is task-dependent —counter-productive on VizDoom Defend-the-Center** (stack 1 beats stack 4 by 2.4 points) even though it is essential in Pong, because the 3D scenario’s optimal play is dominated by spatial information already present in a single frame.

1. Research question and motivation

Research question. *Which algorithmic design choices most affect deep-RL performance, and do those choices behave the same way on a fully-observable 2D game (Pong) as on a partially-observable 3D FPS (VizDoom Defend-the-Center)?*

Deep RL has many moving parts—the algorithm family (value-based vs policy-based), exploration policy, replay buffer, target network, observation representation, action-space granularity—and their interactions depend on the task. A practitioner facing a new image-based control problem has to choose each of these. The literature gives strong defaults for canonical benchmarks (Atari, MuJoCo, ProcGen), but it is less clear how those defaults transfer when, for example, the observation becomes a partial 3D view instead of a side-on 2D playfield. This project examines that transfer.

Two tasks are chosen to span the easy-and-well-studied vs harder-and-perceptual axis:

- **Atari Pong (ALE/Pong-v5)**—the canonical 2D RL benchmark, fully observable from one or a few stacked frames, with a simple ± 1 -per-point reward bounded in $[-21, +21]$.
- **VizDoom Defend-the-Center / Health-Gathering**—first-person 3D scenes where only what the agent currently sees is observable, and the reward structure is sparser and less uniform.

The investigation is decomposed into **five threads**, each isolating one design dimension while keeping the rest fixed:

#	Thread	Aspect probed
T1	Algorithm family	DQN (value-based) vs PPO (policy-based) at equal data
T2	DQN components	Lagged target network, ϵ -greedy decay, replay-buffer size
T3	Partial observability	Frame stacking (1 vs 4) on the 3D task
T4	Action-space design	Discrete vs MultiDiscrete in VizDoom
T5	Task complexity	Final performance across VizDoom scenarios

Each thread is framed so the result yields a one-sentence message. The motivation throughout is informational: the goal is not to set a state-of-the-art score, but to develop *transferable intuition* about which RL design choices matter when, with quantitative evidence on two contrasting tasks.

2. Methods

This section describes the two tasks, the shared image-processing pipeline, the DQN and PPO algorithms, the software stack, and the evaluation protocol. It is intended to be read without referring

to the code; the per-experiment YAML files in `configs/` hold the exact hyperparameter values.

2.1 2.1 Tasks and environments

Two image-based environments are used, chosen to contrast a 2D fully-observable arcade game with a 3D partially-observable first-person shooter:

- **Atari Pong** (ALE/Pong-v5, Arcade Learning Environment [Bellemare et al. 2013; Machado et al. 2018]). A two-paddle game whose reward is the score margin, bounded in $[-21, +21]$. Sticky actions (action-repeat probability 0.25) inject stochasticity so an agent cannot succeed by memorising a fixed action sequence—it must learn a robust policy.
- **VizDoom** [Kempka et al. 2016] via the Farama Gymnasium wrapper, with two scenarios used in the report: **Defend-the-Center** (the headline 3D task) and **Health-Gathering** (a harder scenario used for the difficulty comparison). Both expose a first-person RGB frame and a small vector of game variables; only the frame is fed to the agent.

2.2 2.2 Shared image pipeline

For Atari, Stable-Baselines3’s standard `AtariWrapper` is applied: no-op reset, frame-skip 4, grayscale, resize to 84×84 , reward clipping to $\{-1, 0, +1\}$, and episodic life resetting where applicable. For VizDoom a custom `VizDoomScreenWrapper` mirrors the same shape: it discards the game variables, grayscales the screen, and resizes to 84×84 uint8.

Both pipelines are then wrapped with `VecFrameStack(4)`, so the policy observes a stack of the four most recent frames ($84 \times 84 \times 4$). This provides temporal context—most importantly ball velocity in Pong and projectile/enemy motion in VizDoom. The neural network is the **Nature-CNN** architecture of Mnih et al. [2015]: three convolutional layers (32 / 64 / 64 channels with $8 \times 8 / 4 \times 4 / 3 \times 3$ kernels and strides 4 / 2 / 1) followed by a 512-unit MLP head and the algorithm-specific output layer.

2.3 2.3 DQN (value-based)

DQN [Mnih et al. 2015] approximates the optimal action-value function $Q^*(s, a)$ with a neural network Q_θ . Three components stabilise training of this off-policy bootstrapped target:

- **Experience replay.** Every observed transition $(s, a, r, s', done)$ is stored in a fixed-size buffer D . Each gradient update samples a minibatch from D , decorrelating consecutive samples and bringing the data closer to an i.i.d.-from-the-stationary-distribution assumption.
- **Target network.** A *lagged* copy Q_{θ^-} provides the bootstrap target $y_t = r_t + \gamma \cdot \max_{a'} Q_{\theta^-}(s_{t+1}, a')$. Without it, the regression target moves every gradient step and training is prone to divergence—the so-called “moving-target” problem. θ^- is hard-copied from θ every `target_update_interval` steps.
- **ϵ -greedy exploration.** The behaviour policy picks a uniformly random action with probability ϵ (linearly annealed from $1 \rightarrow 0.01$ over a fraction of training) and the greedy action otherwise.

The loss is the squared TD error $L(\theta) = \mathbb{E}_{\{(s,a,r,s') \sim \mathcal{D}\}}[(y_t - Q_\theta(s,a))^2]$, optimised with Adam, $\gamma = 0.99$. Pong-baseline (P1) settings: replay buffer 500 k, learning starts after 100 k transitions, batch size 32, `train_freq = 4` (one gradient step every four environment steps), target-network sync every 1 000 steps, ϵ annealed over the first 10 % of training to a final 0.01.

2.4 2.4 PPO (policy-based)

PPO [Schulman et al. 2017] directly optimises a stochastic policy $\pi_\theta(a|s)$ in actor-critic form alongside a value baseline $V_\phi(s)$. It collects on-policy rollouts in parallel from many environments, then performs several epochs of minibatch updates on each rollout using the *clipped surrogate objective*

$$L^{CLIP}(\theta) = \mathbb{E}_t [\min(r_t(\theta) A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) A_t)]$$

where $r_t(\theta) = \pi_\theta(a_t|s_t) / \pi_{\theta_{\text{old}}}(a_t|s_t)$ is the importance ratio between the new and old policies. The `clip` prevents updates that move the policy too far from the data-generating distribution, keeping the on-policy approximation valid for several reuse epochs.

Advantages A_t are estimated with **Generalised Advantage Estimation** (GAE, $\lambda = 0.95$), which interpolates between the high-variance Monte-Carlo and the high-bias one-step TD estimators. The full training loss adds a value-function regression term and an entropy bonus to maintain exploration: $L = L^{\text{CLIP}} - c_v \cdot L^V + c_e \cdot H[\pi_\theta]$.

Rollouts are collected from `n_envs = 16` parallel environments (SubprocVecEnv) of length `n_steps = 128`, giving 2 048 transitions per update; each rollout is shuffled into minibatches of 512 for `n_epochs = 4` passes. Other settings: $\gamma = 0.99$, clip range 0.1 (Pong) / 0.2 (VizDoom), entropy coefficient 0.01, value coefficient 0.5, Adam learning rate 2.5×10^{-4} .

2.5 2.5 Software, hardware, and AI-tool disclosure

- **Stable-Baselines3** [Raffin et al. 2021] provides the DQN and PPO implementations and the vectorised-environment plumbing. The `CnnPolicy` selects the Nature-CNN feature extractor described in §2.2.
- **Gymnasium** [Farama Foundation] is the environment API. Atari uses `ale-py` (ROMs bundled); VizDoom uses the `vizdoom.gymnasium_wrapper`.
- **PyTorch** with CUDA 12.4 runs the network forward / backward on the GPU.
- Training was performed on a shared lab workstation (RTX 4090 24 GB, i7-13700K = 24 threads, Ubuntu 24.04) in a Conda environment with Python 3.12. cuDNN autotuning and TF32 matmuls were enabled for throughput.

AI-tool disclosure (course requirement). *Claude Code* (Anthropic) was used to scaffold the repository, implement the training / evaluation / plotting code, build the SSH + scp deployment pipeline, and draft this Methods section and the results-and-figure helpers for the author to review and edit. All experimental design, hyperparameter choices, and the analysis in §4 are the author’s own.

2.6 Experiment protocol

Each (configuration, seed) pair is trained from scratch with a fixed seed. The report uses ≥ 3 **seeds per curve**; learning curves show the mean and the ± 1 standard-deviation band across seeds, computed by interpolating each seed onto a shared timestep grid and aggregating.

At the end of training, every run is evaluated for 20 deterministic episodes (no exploration noise) with a fresh seed; the mean and standard deviation are saved to `eval.json`. During training, `EvalCallback` runs a 10-episode deterministic evaluation every 100 k steps and keeps the best-performing checkpoint, which guards against late-training instability.

To make the cross-algorithm comparison fair, Pong DQN and Pong PPO are both trained for **7 M environment steps**. The VizDoom Defend-Center PPO-vs- $\{\text{discrete/multidiscrete, stack 4/1}\}$ comparisons are at **2 M steps** each. Hyperparameters for every experiment are versioned in `configs/<id>.yaml`.

3. Experiments

Results are organised by the five threads from §1, not by experiment ID, so that each thread tells one story. Every number reported below comes from `results/<exp_id>_s<seed>/eval.json` (mean and standard deviation across the 20-episode deterministic evaluation), aggregated as the mean of seed means $\pm$ its population standard deviation. Training-curve figures show mean $\pm$ 1 SD across seeds.

3.1 T1 —Algorithm family: DQN vs PPO

Aspect. Value-based learning with off-policy replay (DQN) versus on-policy trust-region policy optimisation (PPO), at an *equal environment-step budget*, on **both** environments (Pong and Defend-the-Center). To make the cross-algorithm claim robust on Pong, two PPO hyperparameter settings are compared against DQN.




Figure 1: T1a (left): Pong with DQN, original-recipe PPO (P5) and the SB3-zoo / Schulman 2017 linear-LR-decay recipe (P5b). T1b (right): Defend-the-Center with PPO (V1) and DQN (V5).

Algorithm / variant	Config	Task	Seeds	Mean $\pm$ SD (eval)	Budget
DQN	P1	Pong	3	+4.77 $\pm$ 6.32	7 M
PPO (our recipe)	P5_ppo_pong	Pong	3	-6.58 $\pm$ 1.86	7 M
PPO (SB3-zoo recipe)	P5b_ppo_zoo	Pong	3	-6.07 $\pm$ 3.13	7 M
PPO	V1_defendcenter	Defend-Center	3	+9.37 $\pm$ 0.54	2 M
DQN	V5_dqn_defendcenter	Defend-Center	3	+8.85 $\pm$ 1.19	2 M

Pong (T1a). At an equal 7 M-step budget, value-based DQN substantially outperforms both PPO variants. The original-recipe PPO (constant `learning_rate` = $2.5e-4$, constant `clip_range` =

0.1, batch_size = 512) ends at -6.58 ± 1.86 , well below DQN' s $+4.77 \pm 6.32$. The best DQN seed reaches **+11.15** (clear winning play); the best original-PPO seed stays negative (-3.95).

To control for our PPO config being mistuned, a second variant **P5b_ppo_zoo** uses the SB3-zoo / Schulman 2017 Atari recipe —learning_rate: lin_2.5e-4 (linear decay to 0), clip_range: lin_0.1 (linear decay to 0), batch_size = 256—with everything else identical to P5. P5b ends at -6.07 ± 3.13 , **marginally better on the mean (0.5 pt)** but with **1.7× higher cross-seed standard deviation**. One P5b seed (s1) reaches -1.65 , competitive with DQN; the other two plateau near -8 . *The DQN > PPO inversion at 7 M survives the literature-recipe sanity check*. This is most consistent with **PPO' s known Atari hyperparameter-sensitivity / seed-variance problem** (Henderson et al. 2018, "Deep RL that Matters") rather than a fundamental algorithm gap: PPO can reach the literature numbers, but only with a tuned combination of (recipe × budget × seed) that we do not happen to hit. Mechanistically, DQN' s replay buffer reuses every transition across many off-policy gradient updates, whereas PPO discards each rollout after n_epochs = 4 minibatch passes; at a fixed environment-step budget on Atari, the off-policy method extracts more signal per frame.

Defend-the-Center (T1b). The picture inverts in magnitude but not direction on the 3D scenario at 2 M steps. PPO (V1) reaches $+9.37 \pm 0.54$ kills per episode; DQN (V5) reaches $+8.85 \pm 1.19$. *Both algorithms clearly learn Defend-the-Center*, with PPO marginally ahead (0.52 pt)—well within DQN' s seed- σ of 1.19. V5' s higher cross-seed variance is the familiar value-based seed-instability story; the means are essentially tied.

Takeaway. DQN beats PPO on Pong at 7 M ($+4.77$ vs $-6.07 / -6.58$) robustly across two PPO recipes, driven by replay-buffer sample efficiency. **On Defend-the-Center the two are tied** ($+8.85$ vs $+9.37$), suggesting the 7 M-Pong gap is a sample-efficiency effect that PPO closes once the budget is generous relative to the task' s difficulty.

3.2 3.2 T2 —DQN components on Pong (target network, ϵ -greedy, replay buffer)

Aspect. Which of DQN' s three stabilising components—the lagged target network, the ϵ -greedy schedule, and the size of the replay buffer—matters most for performance and stability on Pong?

T2 is reported at the 2 M-step ablation budget. All four variants below were trained for 2 M steps in the original round-1 pass; the round-2 attempt to re-train them at 7 M was pre-empted by shared-GPU contention before finishing (only P2_targetoff seed 1 completed at $+5.6$, leaving the other seeds with their 2 M evaluations—see §4 limitations). For a clean 3-seed comparison we therefore use the 2 M data. The P1 baseline was additionally extended to 7 M ($+4.77$, see T1) for the algorithm-family comparison. The **P4 small-buffer ablation was run in round-3 catch-up** at the same 2 M budget for a clean comparison against P1' s 2 M snapshot.

Variant	Config	Seeds	Mean $\pm$ SD (eval)	Budget	Δ vs baseline
Baseline	P1	3	-4.17 ± 2.32	2 M	—
Target net OFF	P2_targetoff	3	-13.68 ± 1.95	2 M	-9.51

Variant	Config	Seeds	Mean $\pm$ SD (eval)	Budget	Δ vs baseline
ϵ fast (frac = 0.02)	P3_epsfast	3	-7.73 ± 3.17	2 M	-3.56
ϵ slow (frac = 0.5)	P3_epsslow	3	-7.90 ± 4.78	2 M	-3.73
Small buffer (20 k)	P4_buffersmall	3	-5.03 ± 2.50	2 M	-0.86

The lagged target network has by far the largest effect: removing it (setting `target_update_interval` = 1 so the bootstrap target updates every gradient step) drops mean reward from **-4.17 to -13.68** —a **9.5-point gap**. The effect is robust across seeds —the cross-seed standard deviation is in fact slightly *smaller* without the target net (1.95 vs 2.32), so the gap is a consistent shift in performance, not a single bad seed pulling the mean down. (One 7 M-extended P2 seed reached +5.6, but with only one completed seed at that budget we cannot make a 7 M-vs-7 M claim.)

The ϵ -greedy schedule has a smaller, roughly symmetric effect. Both a fast decay (`exploration_fraction` = 0.02, ϵ at floor by 40 k steps) and a slow decay (`exploration_fraction` = 0.5, floor reached only at 1 M steps) reach **-7.73** and **-7.90** respectively, each costing ≈ 3.6 points relative to the default `exploration_fraction` = 0.1. Too much *or* too little exploration hurts by a comparable amount, so the default is near the sweet spot.

The **replay-buffer size** ablation (P4, 20 k vs P1' s 500 k) reaches **-5.03 ± 2.50** at 2 M —only **0.86 points worse than baseline**, a much milder effect than either removing the target net (-9.5 pt) or mistuning the ϵ schedule (-3.6 pt). Two of the three P4 seeds ($s_1 = -5.40$, $s_2 = -1.80$) land close to or above the P1 mean; the third ($s_0 = -7.90$) is pulled down by training-process contention rather than the small buffer itself (see §4.1). The mechanistic story still holds —a 20 k buffer only keeps the last ≈ 25 episodes of transitions so the i.i.d. assumption behind experience replay is weakly satisfied —but at the 2 M ablation budget the effect is dominated by training noise, not buffer correlation. **The component ordering —target-net ($\gg$) ϵ -schedule ($\approx$) buffer-size —matches the canonical Mnih 2015 / Hessel 2018 Rainbow ablation ordering.**

Takeaway. The lagged target network is by far the most important DQN stabilisation component on Pong (~ 9.5 -pt cost), followed by the ϵ -greedy schedule (~ 3.6 -pt cost, roughly symmetric) and replay-buffer size (~ 0.9 -pt cost). The ordering matches the published ablation literature.

3.3 T3 —Partial observability: frame stacking on VizDoom Defend-Center

Aspect. Does temporal context (4-frame stacking) matter more in a 3D first-person view, where a single frame cannot reveal projectile or enemy motion, than in 2D Pong?




Figure 2: DQN components: P1 baseline, P2 target-net OFF, P3 ϵ -fast / ϵ -slow, P4 small buffer.




Figure 3: Defend-Center: V1 (stack 4) vs V4 (stack 1).

Frame stack	Config	Seeds	Mean $\pm$ SD (eval)	Budget
4 frames	V1 (Discrete)	3	9.37 ± 0.54	2 M
1 frame	V4_stack1	3	11.82 ± 0.88	2 M

Counter to the standard Atari intuition, frame stacking did **not** help on VizDoom Defend-the-Center. With a single 84×84 frame as input, the PPO agent averaged **11.82 ± 0.88** kills per episode, versus **9.37 ± 0.54** for the stack-4 variant —a **+2.4-point** advantage for stack 1. Both variants ran for 2 M steps with three seeds and converged smoothly (small standard deviation, so the comparison is reliable, not noise).

A plausible explanation: Defend-the-Center’s optimal play is predominantly **spatial** —rotate to face the nearest monster, fire, repeat. Monster motion is slow enough that a single frame already exposes the spatial configuration the policy needs to act on. Stacking four frames quadruples the input channel dimension, enlarges what the convolutional encoder must process, and seems to slow learning at this budget without contributing usable temporal signal. This sharply contrasts with Pong, where frame stacking is essentially mandatory because ball velocity (direction *and* speed) is not observable from a single frame.

Takeaway. Frame stacking is **task-dependent**, not universal —it is essential in Pong (where ball velocity is encoded in motion) but actively counter-productive on Defend-the-Center (stack 1 beats stack 4 by 2.4 points), where the dominant signal is already spatial within a single frame.

3.4 T4 —Action-space design on VizDoom Defend-Center

Aspect. The Farama VizDoom wrapper exposes a **Discrete** action space of single-button presses by default (`max_buttons_pressed = 1`). Setting `max_buttons_pressed = 0` gives a **MultiDiscrete** space that lets the agent press any combination of buttons at once. Does the larger combinatorial action space help (more expressive control) or hurt (harder credit assignment at the same budget)?

Action space	Config	Seeds	Mean $\pm$ SD (eval)	Budget
Discrete (one button per step)	V1	3	9.37 ± 0.54	2 M
MultiDiscrete (any button combination)	V2_multibinary	3	6.92 ± 0.94	2 M

The default **Discrete** action space (one button per time-step, `max_buttons_pressed = 1`) substantially out-performed the **MultiDiscrete** alternative (`max_buttons_pressed = 0`, any combination of buttons allowed) at the same 2 M-step budget: **9.37 vs 6.92** mean kills per episode. The MultiDiscrete space is strictly more expressive —it contains every Discrete policy as a special case — but the combinatorial blow-up (2^n joint actions for n available buttons) makes exploration harder.




Figure 4: Defend-Center: V1 (Discrete) vs V2 (MultiDiscrete).

PPO is on-policy, so it cannot reuse exploration well across iterations; at 2 M steps it has not seen enough diverse rollouts to identify which button combinations help.

Takeaway. A smaller, well-chosen action space (Discrete with one button per step) beats a strictly more expressive MultiDiscrete space at equal data on Defend-the-Center, because exploration is harder in the larger combinatorial space.

3.5 T5 —Task complexity: Defend-Center vs Health-Gathering

Aspect. How does the final reward of the same PPO recipe scale with scenario difficulty? Defend-Center is a stationary-shooter scenario; Health-Gathering requires the agent to navigate while collecting items, with sparser reward.




Figure 5: Final eval reward by VizDoom scenario.

Scenario	Config	Seeds	Mean $\pm$ SD (eval)	Budget	Reward scale
Defend-the-Center	V1	3	9.37 $\pm$ 0.54	2 M	+1 per kill ($\approx$ 26 max)
Health-Gathering	V3_healthgathering	3	321.33 $\pm$ 6.44	3 M	accumulated survival + pickup reward

The same PPO recipe (16 parallel envs, 84×84 frame stack, Nature-CNN encoder, clipped objective with GAE) transfers across two qualitatively different VizDoom scenarios. Defend-the-Center, a stationary-shooter task, is solved at **9.37 $\pm$ 0.54** kills per episode after 2 M steps. Health-Gathering, which requires the agent to navigate while picking up health packs to survive, reaches **321.33 $\pm$ 6.44** cumulative reward after 3 M steps. Both have very low cross-seed variance —both scenarios are well within the recipe’s reach when given enough frames.

Direct numerical comparison between the two scenarios is not meaningful because their reward scales differ (single-event kill rewards vs an accumulating survival/pickup signal). The qualitative finding is that the same **architecture and learning algorithm** generalise across both, with the only adjustment being a 50 % larger frame budget for the harder Health-Gathering task.

Takeaway. The same PPO recipe generalises across Defend-the-Center and Health-Gathering —both scenarios are solved with very low cross-seed variance —with the only required adjustment being a 50 % larger frame budget for the harder Health-Gathering task.

3.6 Qualitative examples (gameplay)

Per-run `gameplay.gif` files (under `results/<id>_s0/gameplay.gif`) record one deterministic evaluation episode of each agent. Two are highlighted as the report’s “examples” :

- **results/P1_s0/gameplay.gif** —Pong DQN baseline after 7 M training steps. The paddle reliably tracks the ball, returns most serves, and the final 20-episode evaluation averages a positive score (see T1).
- **results/V1_defendcenter_s0/gameplay.gif** —VizDoom PPO on Defend-the-Center after 2 M training steps. The agent rotates to face incoming demons and fires when on-target, averaging ≈ 9 kills per episode.

4 4. Discussion

The five threads jointly argue that **the most consequential design choices for image-based deep RL are sample-efficiency-related (algorithm family) and stability-related (target network), while observation- and action-representation choices interact with the task in ways that do not always transfer from one benchmark family to another.**

Sample efficiency dominates the Pong algorithm-family comparison. At an equal 7 M-step Pong budget, DQN reached +4.77 mean reward versus $-6.58 / -6.07$ for two distinct PPO recipes (T1). The latter (P5b_ppo_zoo) uses the published SB3-zoo / Schulman 2017 Atari recipe —linear LR decay, linear clip-range decay, `batch_size = 256` —and was added explicitly to test whether the inversion is a hyperparameter artefact. It is not: the zoo recipe improves the mean by only **0.5 pt** while inflating the cross-seed σ from 1.86 to 3.13. One zoo-recipe seed (P5b_sl) lands at -1.65 , competitive with DQN, but the other two plateau near -8 . This pattern is exactly the **PPO-on-Atari seed-variance failure mode** flagged by Henderson et al. 2018 —PPO can reach literature performance, but only within a narrow (recipe $\times$ seed $\times$ budget) sweet spot we did not happen to hit. DQN’s replay buffer reuses every observed transition across many gradient updates, while PPO discards each rollout after four minibatch passes; at a fixed environment-step budget on Atari, the off-policy method extracts more learning signal per frame.

On Defend-the-Center the two algorithms are tied (T1b): PPO +9.37 vs DQN +8.85 at 2 M, a 0.52-point gap that sits inside DQN’s seed- σ of 1.19. The Pong-vs-Defend-Center difference is informative —on a relatively easy 3D scenario both families converge to similar performance well before their respective asymptotic ceilings, so the “DQN beats PPO” claim is **a sample-efficiency story at constrained budget, not an algorithm-family claim.**

Target networks are not optional for DQN. Removing the lagged target network was the single most damaging change in any DQN experiment: P1 baseline -4.17 vs P2-target-off -13.68 at the 2 M ablation budget (T2). The gap of ~ 9.5 points is a consistent shift across all three seeds —the cross-seed standard deviation is actually slightly *tighter* for the target-off variant (1.95 vs 2.32), so the effect is not an instability artefact but a uniform regression. (Of the planned 7 M-budget re-runs of this ablation only one P2 seed completed, at +5.6; with $n = 1$ we cannot make a robust 7 M-vs-7 M claim, so the comparison is anchored at 2 M.) The ϵ -greedy schedule’s effect is smaller and roughly symmetric (≈ 3.6 pt for both faster and slower decays). The **replay-buffer size** matters least at this budget: P4 (20 k buffer) costs only ≈ 0.9 pt vs baseline, and within-ablation seed variance dominates the mean effect. The full ordering —target-net $\gg \epsilon$ -schedule $\approx$ buffer-size —reproduces the canonical Mnih 2015 / Rainbow ablation pattern.

Observation representation is task-dependent, not universal. The most striking single result is T3: frame stacking, which is essential in Pong, *actively hurts* on VizDoom Defend-the-Center (stack 1 at 11.82 vs stack 4 at 9.37). Defend-the-Center’s optimal play is largely instantaneous and spatial —face the nearest monster, fire —so the extra three channels of a stack-4 input slow learning without contributing usable temporal signal. The broader lesson: heuristics that became “default” on one benchmark family (Atari) should not be carried over uncritically to another (3D first-person).

Larger action spaces are not free either. On the same scenario, V1’s Discrete action space

(9.37) cleanly beat V2' s strictly-more-expressive MultiDiscrete space (6.92) at equal data (T4). MultiDiscrete contains the Discrete one as a special case, so the loss is purely an exploration / sample-complexity effect: PPO at 2 M steps has not seen enough rollouts to identify useful joint button combinations, and the wider action distribution makes exploration less efficient. This is T3' s lesson in a different guise: *input/output complexity should match the data budget*, not exceed it.

The PPO recipe generalises across VizDoom scenarios (T5). Both Defend-the-Center (9.37 kills, 2 M) and Health-Gathering (321 reward, 3 M) are solved with very low seed variance, the only difference being budget. Together with T3 / T4 this paints a coherent picture: the algorithm transfers, but the representation choices around it must be re-validated per task.

4.1 4.1 Limitations

Four honest caveats. **First**, the planned round-2 re-training of all DQN ablations at 7 M was preempted by shared-GPU contention (≈ 9 h/run instead of the projected ≈ 3.8 h, plus one seed killed mid-training). Only P1 ran 3-seeds $\times$ 7 M cleanly. P2_targetoff completed 1 of 3 seeds at 7 M (seed 1 at +5.6); seeds 0 and 2 retain their 2 M round-1 evaluations on disk, so the on-disk aggregate (−6.92) mixes budgets and is not a clean 7 M comparison. T2 is therefore reported at the **2 M ablation budget** where all three seeds of every variant are comparable. P1 was additionally extended to 7 M (+4.77) for the T1 algorithm-family comparison, which is a clean 7 M-vs-7 M result against both PPO variants.

Second, the P4_buffersmall seed 0 training-curve CSV is corrupted by a brief concurrent-write incident in round-3: a duplicate queue master was inadvertently launched and ran for ~ 1 minute alongside the intended one, with both python processes writing to the same progress.csv. The final 20-episode eval (run sequentially at the end) is unaffected (P4_s0 = −7.90), so the P4 aggregate in the T2 table is computed from all three eval.json files; only the *training curve* for s0 is dropped from the T2 figure (the figure averages over the surviving s1+s2). The narrative (P4 $\approx -5 \pm 2.5$, 0.9 pt worse than baseline) is robust to this seed' s trajectory being unobservable.

Third, the V5_dqn_defendcenter run is at the 2 M budget (matching V1' s 2 M PPO budget for a fair comparison), but Pong DQN required nearer 7 M to converge. We do not know what V5 would reach at 5–10 M; the $+8.85 \approx V1' s +9.37$ conclusion is specifically for the *small-budget* DC regime.

Fourth, Pong DQN exhibits high cross-seed variance even at 7 M (P1 $\sigma = 6.32$ across three seeds, with one seed regressing to −3.85). PPO shows the inverse pattern —*low* mean but also low σ under our original recipe, and *higher* σ (3.13) under the zoo recipe, with one zoo seed reaching −1.65. A larger seed budget ($n \geq 5$) would tighten central estimates for both algorithms.

4.2 4.2 Future work

1. **Extend Pong PPO to ~ 12 – 15 M steps** to test whether the SB3-zoo recipe —whose LR schedule is tuned for a 10 M+ budget —closes the gap to DQN once the schedule' s decay-to-zero tail is given enough room. The one P5b seed that reached −1.65 suggests this is at least plausible.

2. **Extend V5 (Defend-Center DQN) to 5–10 M** to test whether DQN' s slower sample-efficiency on the 3D scenario eventually exceeds PPO' s 2 M ceiling, mirroring the Pong story at a different budget point.
3. **Add reward normalisation (VecNormalize) to the VizDoom pipeline** so scenarios with large-magnitude or sparse-spike rewards can be trained with the same recipe more reliably.
4. **Run more seeds for the Pong algorithm-family comparison ($n \geq 5$)** to sharpen both the DQN central estimate and the PPO seed-variance picture that the zoo-recipe ablation surfaced.

5. References

- Bellemare, M. G., Naddaf, Y., Veness, J., & Bowling, M. (2013). *The Arcade Learning Environment: An Evaluation Platform for General Agents*. Journal of Artificial Intelligence Research, 47, 253–279. (arXiv:1207.4708)
- Machado, M. C., Bellemare, M. G., Talvitie, E., Veness, J., Hausknecht, M., & Bowling, M. (2018). *Revisiting the Arcade Learning Environment*. Journal of Artificial Intelligence Research, 61, 523–562. (arXiv:1709.06009)
- Kempka, M., Wydmuch, M., Runc, G., Toczek, J., & Jaśkowski, W. (2016). *ViZDoom: A Doom-based AI Research Platform for Visual Reinforcement Learning*. IEEE Conference on Computational Intelligence and Games. (arXiv:1605.02097)
- Mnih, V., Kavukcuoglu, K., Silver, D., et al. (2015). *Human-level Control through Deep Reinforcement Learning*. Nature, 518, 529–533. (doi:10.1038/nature14236)
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). *Proximal Policy Optimization Algorithms*. arXiv:1707.06347.
- Raffin, A., Hill, A., Gleave, A., Kanervisto, A., Ernestus, M., & Dormann, N. (2021). *Stable-Baselines3: Reliable Reinforcement Learning Implementations*. Journal of Machine Learning Research, 22(268), 1–8.
- Farama Foundation. *Gymnasium*. Repository: github.com/Farama-Foundation/Gymnasium.
- Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd ed.). MIT Press.
- Raffin, A. (2020). *RL Baselines3 Zoo*. Repository: github.com/DLR-RM/r1-baselines3-zoo (consulted for Atari hyperparameter defaults).
- **AI tools.** Anthropic *Claude Code* was used as described in §2.5.

6 Appendix A —Repository layout (selected files)

The full project lives at this report's submission alongside the PDF. The most relevant graded files:

```
common/
  train_core.py      config-driven training loop (DQN+PPO)
  envs.py            Atari / VizDoom env factories (shared 84x84 pipeline)
  vizdoom_wrappers.py VizDoom Dict-obs -> single-image wrapper
  callbacks.py       EvalCallback + CheckpointCallback assembly
  eval_utils.py      final evaluation + gameplay-GIF recording
  plotting.py        per-run + cross-seed aggregate curves
  utils.py           config loading, seeding, torch perf flags
pong/
  train_pong_dqn.py  train_pong_ppo.py  eval_record.py
vizdoom/
  train_vizdoom_ppo.py train_vizdoom_dqn.py eval_record.py
configs/
  one YAML per experiment (P1..P5, V0..V5)
tools/
  build_report_assets.py (this report's figures + summary.md)
deploy/
  SSH/scp workflow (server has no git)
```

The full code appendix follows after this page (Appendix B). The same code is also browsable on GitHub at <https://github.com/Ashurali/RLAIHW3>; the in-PDF copy below is a frozen snapshot tied to the commit that produced the results in section 3.

7 Appendix B —Program code

The full source tree is reproduced below, grouped by role for top-down reading. The same files are also available at <https://github.com/Ashurali/RLAIHW3> (branch main); this appendix is a self-contained, frozen snapshot for grading.

Every source file is auto-extracted from the repository by `tools/build_appendix.py`, so what you read here is exactly what trained the models behind the results in §3. The code is documented inline —each module opens with a docstring explaining its role, and non-obvious decisions (schedule resolver, frame-stack wrapper, CSV-corruption fallback, GPU perf flags) carry comments at point-of-use.

7.1 B.1 Common engine (algorithm-agnostic training & evaluation stack)

7.1.1 common/utils.py (101 lines)

```

"""Configuration loading, seeding and run-directory helpers."""
from __future__ import annotations

import random
from pathlib import Path

import numpy as np
import yaml

# common/ lives one level below the repository root.
REPO_ROOT = Path(__file__).resolve().parents[1]
RESULTS_DIR = REPO_ROOT / "results"

def load_config(path) -> dict:
    """Load a YAML experiment config into a plain dict."""
    with open(path, "r", encoding="utf-8") as f:
        cfg = yaml.safe_load(f)
    for required in ("exp_id", "task", "algo", "env_id", "total_timesteps"):
        if required not in cfg:
            raise ValueError(f"Config {path} is missing required key '{required}'.")
    return cfg

def _linear_schedule(initial_value: float):
    """SB3 schedule: linear from ``initial_value`` (progress=1) down to 0.
    SB3 invokes the callable with ``progress_remaining`` going 1.0 -> 0.0 over
    training. This is the standard schedule used by the SB3-zoo / Schulman
    PPO-Atari recipe: linear LR-decay and clip-range-decay both critical for
    Pong convergence.
    """
    def schedule(progress_remaining: float) -> float:
        return float(progress_remaining) * float(initial_value)
    schedule.__name__ = f"lin_{initial_value:g}"
    return schedule

def resolve_schedules(kwarg: dict) -> dict:
    """Convert ``lin_{float}>`` strings in algo_kwarg into SB3 schedule callables.

```

Lets configs say ``learning\_rate: lin\_2.5e-4`` and have SB3 receive a proper linear-decay callable. Returns a NEW dict (does not mutate input).

```
"""
out = dict(kwargs)
for k, v in list(out.items()):
    if isinstance(v, str) and v.startswith("lin_"):
        try:
            init = float(v[len("lin_"):])
        except ValueError:
            continue
        out[k] = _linear_schedule(init)
return out
```

```
def set_global_seeds(seed: int) -> None:
    """Seed Python, NumPy and (when present) PyTorch for reproducibility."""
    random.seed(seed)
    np.random.seed(seed)
    try:
        import torch

        torch.manual_seed(seed)
        if torch.cuda.is_available():
            torch.cuda.manual_seed_all(seed)
    except ImportError:
        # torch is only required on the training server, not for tooling.
        pass
```

```
def configure_torch_perf() -> None:
    """Enable throughput-oriented GPU settings (safe on any CUDA GPU).

    cuDNN autotuning picks the fastest conv algorithms for our fixed 84x84
    inputs; TF32 matmuls use the Ampere/Ada tensor cores. Both trade a little
    numerical determinism for speed, which is fine for RL throughput.
    """
    try:
        import torch

        torch.backends.cudnn.benchmark = True
        if torch.cuda.is_available():
            torch.backends.cuda.matmul.allow_tf32 = True
            torch.backends.cudnn.allow_tf32 = True
            torch.set_float32_matmul_precision("high")
    except ImportError:
```

pass

```
def get_run_dir(exp_id: str, seed: int, create: bool = True) -> Path:
    """Return ``results/<exp_id>_s<seed>/``; create it unless told otherwise."""
    run_dir = RESULTS_DIR / f"{exp_id}_s{seed}"
    if create:
        run_dir.mkdir(parents=True, exist_ok=True)
    return run_dir

def save_config_copy(cfg: dict, run_dir: Path) -> None:
    """Persist the exact config used for a run next to its outputs."""
    with open(Path(run_dir) / "config.yaml", "w", encoding="utf-8") as f:
        yaml.safe_dump(cfg, f, sort_keys=False)
```

7.1.2 common/envs.py (95 lines)

"""Environment factories for both tasks, with a unified image front-end.

Both Pong and VizDoom are turned into a ``VecFrameStack``'d image ``VecEnv`` so the exact same ``CnnPolicy`` code trains on each. :func:`build\_env\_fn` returns a closure ``env\_fn(seed, n\_envs, eval\_mode, render\_mode)`` that the training core uses for the training, evaluation and rendering environments alike.

```
from __future__ import annotations

from stable_baselines3.common.env_util import make_atari_env, make_vec_env
from stable_baselines3.common.vec_env import (
    DummyVecEnv,
    SubprocVecEnv,
    VecFrameStack,
)

import gymnasium as gym
import ale_py
gym.register_envs(ale_py)

from common.vizdoom_wrappers import VizDoomScreenWrapper

def make_atari_vec(env_id, n_envs, seed, n_stack=4, render_mode=None):
    """Atari pipeline: AtariWrapper (grayscale/resize/skip) + frame stack."""
    env_kwargs = {"render_mode": render_mode} if render_mode else None
    venv = make_atari_env(env_id, n_envs=n_envs, seed=seed, env_kwargs=env_kwargs)
```

```

    return VecFrameStack(venv, n_stack=n_stack)

def make_vizdoom_vec(
    env_id,
    n_envs,
    seed,
    n_stack=4,
    frame_skip=4,
    obs_shape=(84, 84),
    grayscale=True,
    env_kwargs=None,
    render_mode=None,
):
    """VizDoom pipeline mirroring Atari: screen -> 84x84 (+gray) + frame stack."""
    kwargs = dict(env_kwargs or {})
    kwargs.setdefault("frame_skip", frame_skip)
    if render_mode:
        kwargs["render_mode"] = render_mode

    # SubprocVecEnv parallelises env stepping (the CPU bottleneck) for PPO;
    # a single env (DQN / evaluation) stays in-process for simplicity.
    vec_cls = SubprocVecEnv if n_envs > 1 else DummyVecEnv
    venv = make_vec_env(
        env_id,
        n_envs=n_envs,
        seed=seed,
        env_kwargs=kwargs,
        wrapper_class=VizDoomScreenWrapper,
        wrapper_kwargs={"shape": tuple(obs_shape), "grayscale": grayscale},
        vec_env_cls=vec_cls,
    )
    return VecFrameStack(venv, n_stack=n_stack)

def build_env_fn(cfg):
    """Return ``env_fn(seed, n_envs, eval_mode, render_mode)`` for ``cfg['task']``."""
    task = cfg["task"]

    if task == "pong":
        def env_fn(seed, n_envs, eval_mode=False, render_mode=None):
            return make_atari_vec(
                cfg["env_id"],
                n_envs=1 if eval_mode else n_envs,
                seed=seed,

```

```

        n_stack=cfg.get("n_stack", 4),
        render_mode=render_mode,
    )

    return env_fn

if task == "vizdoom":
    def env_fn(seed, n_envs, eval_mode=False, render_mode=None):
        return make_vizdoom_vec(
            cfg["env_id"],
            n_envs=1 if eval_mode else n_envs,
            seed=seed,
            n_stack=cfg.get("n_stack", 4),
            frame_skip=cfg.get("frame_skip", 4),
            obs_shape=cfg.get("obs_shape", (84, 84)),
            grayscale=cfg.get("grayscale", True),
            env_kwargs=cfg.get("env_kwargs"),
            render_mode=render_mode,
        )

    return env_fn

    raise ValueError(f"Unknown task '{task}' (expected 'pong' or 'vizdoom').")

```

7.1.3 common/vizdoom_wrappers.py (63 lines)

"""VizDoom -> SB3 image pipeline.

*The Farama Gymnasium wrapper for ViZDoom returns a *Dict* observation: a ``screen`` image plus a ``gamevariables`` vector. SB3's ``CnnPolicy`` expects a single image tensor, so this wrapper drops the game variables and converts the screen to the same 84x84 (optionally grayscale) uint8 format used for Atari.*

With a subsequent ``VecFrameStack(4)`` this reproduces the classic Atari DQN/PPO front-end, giving Pong and VizDoom an identical observation pipeline so the value-based vs policy-based comparison is not confounded by preprocessing.

```
from __future__ import annotations
```

```
import cv2
import gymnasium as gym
import numpy as np
from gymnasium import spaces
```

Importing the wrapper module registers all "Vizdoom-v*" Gymnasium ids as a*

```
# side effect. Requires vizdoom >= 1.2 (Gymnasium support).
import vizdoom.gymnasium_wrapper # noqa: F401
```

```
class VizDoomScreenWrapper(gym.ObservationWrapper):
```

```
    """Extract the screen buffer, resize it and optionally grayscale it.
```

```
    Output observation: a uint8 array of shape ``(H, W, C)`` (channel-last,  

``C == 1`` when grayscale else ``3``). Channel-last is what SB3's  

``VecTransposeImage`` expects before it auto-converts to channel-first.  

"""
```

```
    def __init__(self, env, shape=(84, 84), grayscale: bool = True):
```

```
        super().__init__(env)
        self.shape = tuple(shape)
        self.grayscale = grayscale
        channels = 1 if grayscale else 3
        self.observation_space = spaces.Box(
            low=0,
            high=255,
            shape=(self.shape[0], self.shape[1], channels),
            dtype=np.uint8,
        )
```

```
    def observation(self, obs):
```

```
        # The Dict obs carries the frame under "screen"; tolerate a bare array.
        screen = obs["screen"] if isinstance(obs, dict) else obs
        screen = np.asarray(screen)
```

```
        # Normalize to channel-last (H, W, C); some screen formats are (C, H, W).
```

```
        if screen.ndim == 3 and screen.shape[0] in (1, 3) and screen.shape[2] not in (1, 3):
            screen = np.transpose(screen, (1, 2, 0))
```

```
        if self.grayscale:
```

```
            screen = cv2.cvtColor(screen, cv2.COLOR_RGB2GRAY)
```

```
        # cv2.resize takes (width, height); self.shape is (height, width).
```

```
        screen = cv2.resize(
            screen, (self.shape[1], self.shape[0]), interpolation=cv2.INTER_AREA
        )
```

```
        if self.grayscale:
```

```
            screen = screen[:, :, None]
```

```
        return screen.astype(np.uint8)
```

7.1.4 common/callbacks.py (34 lines)

```

"""Callback assembly: periodic deterministic evaluation + checkpointing."""
from __future__ import annotations

from pathlib import Path

from stable_baselines3.common.callbacks import CheckpointCallback, EvalCallback

def make_callbacks(cfg, run_dir: Path, eval_env):
    """Build an ``EvalCallback`` (best model + eval log) and a checkpointer.

    Frequencies in the config are expressed in *environment steps*. SB3 counts
    callback ticks per environment, so we divide by ``n_envs`` to keep the
    wall-clock cadence stable regardless of how many parallel envs a run uses.
    """

    n_envs = max(int(cfg.get("n_envs", 1)), 1)
    eval_freq = max(int(cfg.get("eval_freq", 100_000)) // n_envs, 1)
    ckpt_freq = max(int(cfg.get("checkpoint_freq", 250_000)) // n_envs, 1)

    eval_cb = EvalCallback(
        eval_env,
        best_model_save_path=str(Path(run_dir) / "best_model"),
        log_path=str(run_dir),
        eval_freq=eval_freq,
        n_eval_episodes=int(cfg.get("eval_episodes", 10)),
        deterministic=True,
        render=False,
    )
    ckpt_cb = CheckpointCallback(
        save_freq=ckpt_freq,
        save_path=str(Path(run_dir) / "checkpoints"),
        name_prefix=cfg["exp_id"],
    )
    return [eval_cb, ckpt_cb]

```

7.1.5 common/eval_utils.py (124 lines)

```

"""Final evaluation and gameplay-GIF recording.

Shared by the training core (so every run auto-produces ``eval.json`` and
``gameplay.gif``) and by the standalone ``eval_record.py`` entry points (so a
finished model can be re-evaluated / re-recorded without retraining).
    """

```

```
from __future__ import annotations

import json
from pathlib import Path

import imageio
import numpy as np
from stable_baselines3.common.evaluation import evaluate_policy


def _base_envs(venv):
    """Drill through VecEnvWrappers (e.g. VecFrameStack) to the env list."""
    while hasattr(venv, "venv"):
        venv = venv.venv
    return getattr(venv, "envs", None)


def evaluate_and_record(
    model,
    env_fn,
    cfg,
    run_dir: Path,
    seed: int,
    record_gif: bool = True,
    gif_episodes: int = 1,
    max_gif_steps: int = 3000,
):
    """Evaluate the policy and optionally record a gameplay GIF.

    Writes ``eval.json`` (mean/std reward over ``cfg['eval_episodes']``) and,
    when ``record_gif``, ``gameplay.gif``. Returns the stats dict.
    """

    run_dir = Path(run_dir)

    eval_env = env_fn(seed=seed + 10_000, n_envs=1, eval_mode=True)
    mean_r, std_r = evaluate_policy(
        model,
        eval_env,
        n_eval_episodes=int(cfg.get("eval_episodes", 20)),
        deterministic=True,
    )
    eval_env.close()

    stats = {
        "exp_id": cfg["exp_id"],
```



```

        "seed": seed,
        "n_eval_episodes": int(cfg.get("eval_episodes", 20)),
        "mean_reward": float(mean_r),
        "std_reward": float(std_r),
    }
    with open(run_dir / "eval.json", "w", encoding="utf-8") as f:
        json.dump(stats, f, indent=2)
    print(f"[eval] {cfg['exp_id']} seed={seed}: "
          f"{mean_r:.2f} +/- {std_r:.2f} over {stats['n_eval_episodes']} eps")

    if record_gif:
        # Rendering is best-effort: a recording failure must never fail a run.
        try:
            _record_gif(model, env_fn, seed, run_dir, gif_episodes, max_gif_steps)
        except Exception as exc: # noqa: BLE001
            print(f"[eval] GIF recording skipped: {exc}")

    return stats

def _record_gif(model, env_fn, seed, run_dir, gif_episodes, max_gif_steps):
    render_env = env_fn(
        seed=seed + 20_000, n_envs=1, eval_mode=True, render_mode="rgb_array"
    )
    base = _base_envs(render_env)
    frames = []
    obs = render_env.reset()
    episodes, steps = 0, 0
    while episodes < gif_episodes and steps < max_gif_steps:
        action, _ = model.predict(obs, deterministic=True)
        obs, _, done, _ = render_env.step(action)
        if base is not None:
            frame = base[0].render()
            if frame is not None:
                frames.append(np.asarray(frame))
            steps += 1
        if bool(done[0]):
            episodes += 1
    render_env.close()

    if frames:
        imageio.mimsave(run_dir / "gameplay.gif", frames, fps=30)
        print(f"[eval] saved {len(frames)} frames -> {run_dir / 'gameplay.gif'}")

```

```
def run_eval_cli():
    """CLI shared by ``pong/eval_record.py`` and ``vizdoom/eval_record.py``."""
    import argparse

    from stable_baselines3 import DQN, PPO

    from common.envs import build_env_fn
    from common.utils import get_run_dir, load_config

    algos = {"dqn": DQN, "ppo": PPO}

    ap = argparse.ArgumentParser(description="Re-evaluate and record a trained run.")
    ap.add_argument("--config", required=True, help="path to the run's YAML config")
    ap.add_argument("--seed", type=int, default=None, help="override config seed")
    ap.add_argument("--no-gif", action="store_true", help="skip GIF recording")
    args = ap.parse_args()

    cfg = load_config(args.config)
    if args.seed is not None:
        cfg["seed"] = args.seed
    seed = int(cfg.get("seed", 0))

    run_dir = get_run_dir(cfg["exp_id"], seed, create=False)
    model = algos[cfg["algo"]].load(str(run_dir / "model"))
    env_fn = build_env_fn(cfg)
    evaluate_and_record(
        model, env_fn, cfg, run_dir, seed, record_gif=not args.no_gif
    )
```

7.1.6 common/plotting.py (133 lines)

"""Plot per-run learning curves and aggregate seeds into report figures.

*Per the plan, every figure in the report must be built here from data under
``results/`` only -- never hand-edited. Outputs land in ``report\_assets/``.
"""*

```
from __future__ import annotations

import argparse
import glob
from pathlib import Path

import matplotlib

matplotlib.use("Agg") # headless-safe on the training server
```

```

import matplotlib.pyplot as plt # noqa: E402
import numpy as np # noqa: E402
import pandas as pd # noqa: E402

from common.utils import REPO_ROOT, RESULTS_DIR # noqa: E402

X_COL = "time/total_timesteps"
Y_COL = "rollout/ep_rew_mean"
ASSETS_DIR = REPO_ROOT / "report_assets"

def _load_metrics(run_dir):
    """Return (timesteps, ep_rew_mean) arrays from a run's metrics CSV.

    Coerces both columns to numeric and drops non-numeric rows. This handles
    edge cases like concurrent-write corruption where a column might contain
    string garbage instead of floats; we keep whatever valid rows survive.
    Raises ValueError if no usable rows remain.
    """
    run_dir = Path(run_dir)
    path = run_dir / "metrics.csv"
    if not path.exists():
        path = run_dir / "progress.csv"
    df = pd.read_csv(path)
    df = df[[X_COL, Y_COL]].copy()
    df[X_COL] = pd.to_numeric(df[X_COL], errors="coerce")
    df[Y_COL] = pd.to_numeric(df[Y_COL], errors="coerce")
    df = df.dropna()
    if df.empty:
        raise ValueError(f"No usable rows in {path}")
    return df[X_COL].to_numpy(dtype=float), df[Y_COL].to_numpy(dtype=float)

def plot_curve(run_dir):
    """Single-run training curve -> ``<run_dir>/curve.png``."""
    x, y = _load_metrics(run_dir)
    fig, ax = plt.subplots(figsize=(6, 4))
    ax.plot(x, y, color="tab:blue")
    ax.set_xlabel("Timesteps")
    ax.set_ylabel("Episode reward (mean)")
    ax.set_title(Path(run_dir).name)
    ax.grid(alpha=0.3)
    fig.tight_layout()
    out = Path(run_dir) / "curve.png"
    fig.savefig(out, dpi=120)

```

```
plt.close(fig)
return out
```

```
def aggregate_seeds(exp_id, n_points=200):
    """Mean +/- std curve across ``results/<exp_id>_s*/`` runs.

    Curves are interpolated onto a shared timestep grid (the overlap of all
    seeds) so they can be averaged even when logging cadence differs. Runs
    with unparseable metrics CSVs (e.g. concurrent-write corruption) are
    skipped with a printed note; the aggregate proceeds across surviving
    seeds. Raises FileNotFoundError if NO seeds are usable.
    """
    run_dirs = sorted(glob.glob(str(RESULTS_DIR / f"{exp_id}_s*")))
    if not run_dirs:
        raise FileNotFoundError(f"No runs found for exp_id '{exp_id}'.")
    curves = []
    for d in run_dirs:
        try:
            curves.append(_load_metrics(d))
        except (ValueError, KeyError) as exc:
            print(f" [aggregate_seeds] skip {Path(d).name}: {exc}")
    if not curves:
        raise FileNotFoundError(f"No usable metrics CSVs for exp_id '{exp_id}'.")
    x_min = max(c[0].min() for c in curves)
    x_max = min(c[0].max() for c in curves)
    grid = np.linspace(x_min, x_max, n_points)
    stacked = np.stack([np.interp(grid, x, y) for x, y in curves])
    return grid, stacked.mean(axis=0), stacked.std(axis=0)

def plot_comparison(exp_ids, labels=None, title="", out_name="comparison.png"):
    """Overlay mean +/- std curves for several experiments -> report_assets/."""
    labels = labels or exp_ids
    fig, ax = plt.subplots(figsize=(7, 4.5))
    for exp_id, label in zip(exp_ids, labels):
        grid, mean, std = aggregate_seeds(exp_id)
        ax.plot(grid, mean, label=label)
        ax.fill_between(grid, mean - std, mean + std, alpha=0.2)
    ax.set_xlabel("Timesteps")
    ax.set_ylabel("Episode reward (mean)")
    ax.set_title(title)
    ax.legend()
    ax.grid(alpha=0.3)
    fig.tight_layout()
```

```

    ASSETS_DIR.mkdir(parents=True, exist_ok=True)
    out = ASSETS_DIR / out_name
    fig.savefig(out, dpi=150)
    plt.close(fig)
    return out

def main():
    parser = argparse.ArgumentParser(description="Plot RL training curves.")
    sub = parser.add_subparsers(dest="cmd", required=True)

    pc = sub.add_parser("curve", help="single-run curve from a run directory")
    pc.add_argument("run_dir")

    cmp = sub.add_parser("compare", help="overlay several experiments (seeds avg)")
    cmp.add_argument("--ids", nargs="+", required=True)
    cmp.add_argument("--labels", nargs="+")
    cmp.add_argument("--title", default="")
    cmp.add_argument("--out", default="comparison.png")

    args = parser.parse_args()
    if args.cmd == "curve":
        print(plot_curve(args.run_dir))
    elif args.cmd == "compare":
        print(plot_comparison(args.ids, args.labels, args.title, args.out))

if __name__ == "__main__":
    main()

```

7.1.7 common/train_core.py (82 lines)

```

"""Config-driven training loop shared by all four (task, algo) entry points."""
from __future__ import annotations

import shutil
from pathlib import Path

from stable_baselines3.common.logger import configure

from common.callbacks import make_callbacks
from common.envs import build_env_fn
from common.eval_utils import evaluate_and_record
from common.plotting import plot_curve
from common.utils import (

```

```

    configure_torch_perf,
    get_run_dir,
    resolve_schedules,
    save_config_copy,
    set_global_seeds,
)

def run_training(cfg, model_cls):
    """Train one run end-to-end and emit the standard artifact set.

    Under ``results/<exp_id>_s<seed>/`` this produces: ``config.yaml``,
    ``metrics.csv``, ``curve.png``, ``eval.json``, ``gameplay.gif``,
    ``model.zip``, plus ``best_model/`` and ``checkpoints/``.
    """
    seed = int(cfg.get("seed", 0))
    set_global_seeds(seed)
    configure_torch_perf()

    run_dir = get_run_dir(cfg["exp_id"], seed)
    save_config_copy(cfg, run_dir)

    env_fn = build_env_fn(cfg)
    n_envs = int(cfg.get("n_envs", 1))
    train_env = env_fn(seed=seed, n_envs=n_envs, eval_mode=False)
    eval_env = env_fn(seed=seed + 1000, n_envs=1, eval_mode=True)

    # Translate `lin_<float>` strings (e.g. `learning_rate: lin_2.5e-4`) into
    # SB3 linear-decay callables. Required for the literature-recipe PPO config.
    algo_kwargs = resolve_schedules(cfg.get("algo_kwargs", {})) or {}

    model = model_cls(
        cfg.get("policy", "CnnPolicy"),
        train_env,
        seed=seed,
        device=cfg.get("device", "cuda"),
        verbose=1,
        tensorboard_log=str(run_dir / "tb"),
        **algo_kwargs,
    )
    # Mirror stdout into a CSV (-> progress.csv) and TensorBoard event files.
    model.set_logger(configure(str(run_dir), ["stdout", "csv", "tensorboard"]))

    callbacks = make_callbacks(cfg, run_dir, eval_env)
    model.learn(

```

```

        total_timesteps=int(cfg["total_timesteps"]),
        callback=callbacks,
        progress_bar=True,
    )

    model.save(str(run_dir / "model"))
    train_env.close()
    eval_env.close()

    # The plan's per-run deliverable is metrics.csv; SB3 emits progress.csv.
    progress = run_dir / "progress.csv"
    if progress.exists():
        shutil.copyfile(progress, run_dir / "metrics.csv")

    evaluate_and_record(
        model, env_fn, cfg, run_dir, seed, record_gif=cfg.get("record_gif", True)
    )
    try:
        plot_curve(run_dir)
    except Exception as exc: # noqa: BLE001
        print(f"[train_core] curve plotting skipped: {exc}")

    print(f"[train_core] done -> {run_dir}")
    return run_dir

```

7.2 B.2 Per-task training & evaluation entry points

7.2.1 pong/train_pong_dqn.py (37 lines)

"""Train DQN on Atari Pong (Task 1 -- value-based baseline and ablations).

Examples:

python pong/train_pong_dqn.py --config configs/P1.yaml --seed 0

python pong/train_pong_dqn.py --config configs/P2_targetoff.yaml --seed 1

"""

```
from __future__ import annotations
```

```
import argparse
```

```
import sys
```

```
from pathlib import Path
```

Make the repo root importable so "common" resolves when run as a script.

```
sys.path.insert(0, str(Path(__file__).resolve().parents[1]))
```

```
from stable_baselines3 import DQN # noqa: E402
```

```
from common.train_core import run_training # noqa: E402
```

```
from common.utils import load_config # noqa: E402
```

```
def main():
```

```
    ap = argparse.ArgumentParser(description="Train DQN on Atari Pong.")
```

```
    ap.add_argument("--config", required=True)
```

```
    ap.add_argument("--seed", type=int, default=None,
```

```
                    help="override config seed (use 0/1/2 for >=3-seed curves)")
```

```
    args = ap.parse_args()
```

```
    cfg = load_config(args.config)
```

```
    if args.seed is not None:
```

```
        cfg["seed"] = args.seed
```

```
    assert cfg["algo"] == "dqn", f"expected algo: dqn, got {cfg['algo']}"
```

```
    run_training(cfg, DQN)
```

```
if __name__ == "__main__":
```

```
    main()
```

7.2.2 pong/train_pong_ppo.py (36 lines)

"""Train PPO on Atari Pong (Task 1 -- policy-based arm of the P5 comparison).

Example:

```
python pong/train_pong_ppo.py --config configs/P5_ppo_pong.yaml --seed 0
"""
from __future__ import annotations

import argparse
import sys
from pathlib import Path

# Make the repo root importable so "common" resolves when run as a script.
sys.path.insert(0, str(Path(__file__).resolve().parents[1]))

from stable_baselines3 import PPO # noqa: E402

from common.train_core import run_training # noqa: E402
from common.utils import load_config # noqa: E402

def main():
    ap = argparse.ArgumentParser(description="Train PPO on Atari Pong.")
    ap.add_argument("--config", required=True)
    ap.add_argument("--seed", type=int, default=None,
                    help="override config seed (use 0/1/2 for >=3-seed curves)")
    args = ap.parse_args()

    cfg = load_config(args.config)
    if args.seed is not None:
        cfg["seed"] = args.seed
    assert cfg["algo"] == "ppo", f"expected algo: ppo, got {cfg['algo']}"
    run_training(cfg, PPO)

if __name__ == "__main__":
    main()
```

7.2.3 vizdoom/train_vizdoom_dqn.py (39 lines)

"""Train DQN on a VizDoom scenario (Task 2 -- value-based arm of V5).

DQN needs a Discrete action space, so keep the default action config
(max_buttons_pressed = 1) for this run rather than the MultiDiscrete variant.

Example:

```
python vizdoom/train_vizdoom_dqn.py --config configs/V5_dqn_defendcenter.yaml --seed 0
"""
```

```

from __future__ import annotations

import argparse
import sys
from pathlib import Path

# Make the repo root importable so "common" resolves when run as a script.
sys.path.insert(0, str(Path(__file__).resolve().parents[1]))

from stable_baselines3 import DQN # noqa: E402

from common.train_core import run_training # noqa: E402
from common.utils import load_config # noqa: E402

def main():
    ap = argparse.ArgumentParser(description="Train DQN on VizDoom.")
    ap.add_argument("--config", required=True)
    ap.add_argument("--seed", type=int, default=None,
                    help="override config seed (use 0/1/2 for >=3-seed curves)")
    args = ap.parse_args()

    cfg = load_config(args.config)
    if args.seed is not None:
        cfg["seed"] = args.seed
    assert cfg["algo"] == "dqn", f"expected algo: dqn, got {cfg['algo']}"
    run_training(cfg, DQN)

if __name__ == "__main__":
    main()

```

7.2.4 vizdoom/train_vizdoom_ppo.py (39 lines)

"""Train PPO on a VizDoom scenario (Task 2 -- policy-based primary)."""

Examples:

```

python vizdoom/train_vizdoom_ppo.py --config configs/V0_basic.yaml --seed 0
python vizdoom/train_vizdoom_ppo.py --config configs/V1_defendcenter.yaml --seed 0
"""

from __future__ import annotations

import argparse
import sys
from pathlib import Path

```

```

# Make the repo root importable so "common" resolves when run as a script.
# (This folder is intentionally NOT a package, so it cannot shadow the
# installed "vizdoom" pip package that common.vizdoom_wrappers imports.)
sys.path.insert(0, str(Path(__file__).resolve().parents[1]))

from stable_baselines3 import PPO # noqa: E402

from common.train_core import run_training # noqa: E402
from common.utils import load_config # noqa: E402

def main():
    ap = argparse.ArgumentParser(description="Train PPO on VizDoom.")
    ap.add_argument("--config", required=True)
    ap.add_argument("--seed", type=int, default=None,
                    help="override config seed (use 0/1/2 for >=3-seed curves)")
    args = ap.parse_args()

    cfg = load_config(args.config)
    if args.seed is not None:
        cfg["seed"] = args.seed
    assert cfg["algo"] == "ppo", f"expected algo: ppo, got {cfg['algo']}"
    run_training(cfg, PPO)

if __name__ == "__main__":
    main()

```

7.2.5 pong/eval_record.py (17 lines)

"""Re-evaluate a trained Pong model and (re)record its gameplay GIF.

Example:

```
python pong/eval_record.py --config configs/P1.yaml --seed 0
"""
```

```

from __future__ import annotations

import sys
from pathlib import Path

# Make the repo root importable so "common" resolves when run as a script.
sys.path.insert(0, str(Path(__file__).resolve().parents[1]))

from common.eval_utils import run_eval_cli # noqa: E402

```

```
if __name__ == "__main__":  
    run_eval_cli()
```

7.2.6 vizdoom/eval_record.py (17 lines)

"""Re-evaluate a trained VizDoom model and (re)record its gameplay GIF.

Example:

python vizdoom/eval_record.py --config configs/V1_defendcenter.yaml --seed 0
"""

```
from __future__ import annotations
```

```
import sys
```

```
from pathlib import Path
```

Make the repo root importable so "common" resolves when run as a script.

```
sys.path.insert(0, str(Path(__file__).resolve().parents[1]))
```

```
from common.eval_utils import run_eval_cli # noqa: E402
```

```
if __name__ == "__main__":  
    run_eval_cli()
```

7.3 B.3 Experiment configs (one YAML per (task, algorithm, ablation))

7.3.1 configs/P1.yaml (34 lines)

```
# P1 - Pong DQN baseline (Tier A).
# Reference run. Also serves as the reused arm for several ablations:
#   P2 = "target net ON", P3 = "default epsilon", P4 = "large buffer".
# Run >=3 seeds for the report: --seed 0 / 1 / 2
exp_id: P1
task: pong
algo: dqn
env_id: ALE/Pong-v5
policy: CnnPolicy
n_envs: 1
n_stack: 4
seed: 0
total_timesteps: 7000000      # ~3.8 h on the 4090; Long enough to plateau (~+20).
                                # 2M left it undertrained (curve still rising at -3).

eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 1.0e-4
  buffer_size: 500000          # "Large" buffer (~28GB RAM, frame-stacked uint8;
                                # pre-allocated, not dynamic). Fits the shared server
                                # (~42GB free). Raise to 1000000 only if dedicated.
                                # P4 uses a small buffer for the ablation.

  learning_starts: 100000
  batch_size: 32
  gamma: 0.99
  train_freq: 4
  gradient_steps: 1
  target_update_interval: 1000 # lagged target network ON (P2 turns this off)
  exploration_fraction: 0.1    # default epsilon schedule (P3 varies this)
  exploration_initial_eps: 1.0
  exploration_final_eps: 0.01
```

7.3.2 configs/P2_targetoff.yaml (31 lines)

```
# P2 - DQN with the lagged target network effectively OFF (Tier A).
# Aspect: stability. Setting target_update_interval = 1 syncs the target net to
# the online net every gradient step, removing the lag that tames the
# moving-target problem. Compare against P1 (target_update_interval = 1000).
# Everything else is identical to P1.
```

```

exp_id: P2_targetoff
task: pong
algo: dqn
env_id: ALE/Pong-v5
policy: CnnPolicy
n_envs: 1
n_stack: 4
seed: 0
total_timesteps: 7000000      # match P1 (extended from 2M)
eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 1.0e-4
  buffer_size: 500000          # match P1 (~28GB; sized for shared server RAM)
  learning_starts: 100000
  batch_size: 32
  gamma: 0.99
  train_freq: 4
  gradient_steps: 1
  target_update_interval: 1    # <-- target net lag removed (the ablation)
  exploration_fraction: 0.1
  exploration_initial_eps: 1.0
  exploration_final_eps: 0.01

```

7.3.3 configs/P3_epsfast.yaml (30 lines)

```

# P3a - DQN with FAST epsilon decay (Tier A).
# Aspect: exploration vs exploitation. Epsilon reaches its floor very early
# (exploration_fraction = 0.02), so the agent commits to exploitation quickly.
# Compare against P1 (0.1, default) and P3_epsslow (0.5, slow).
exp_id: P3_epsfast
task: pong
algo: dqn
env_id: ALE/Pong-v5
policy: CnnPolicy
n_envs: 1
n_stack: 4
seed: 0
total_timesteps: 7000000      # match P1 (extended from 2M)
eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000

```

```

device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 1.0e-4
  buffer_size: 500000          # match P1 (~28GB; sized for shared server RAM)
  learning_starts: 100000
  batch_size: 32
  gamma: 0.99
  train_freq: 4
  gradient_steps: 1
  target_update_interval: 1000
  exploration_fraction: 0.02    # <-- fast decay (the ablation)
  exploration_initial_eps: 1.0
  exploration_final_eps: 0.01

```

7.3.4 configs/P3_epsslow.yaml (30 lines)

```

# P3b - DQN with SLOW epsilon decay (Tier A).
# Aspect: exploration vs exploitation. Epsilon decays over half of training
# (exploration_fraction = 0.5), keeping the agent exploratory for much longer.
# Compare against P1 (0.1, default) and P3_epsfast (0.02, fast).
exp_id: P3_epsslow
task: pong
algo: dqn
env_id: ALE/Pong-v5
policy: CnnPolicy
n_envs: 1
n_stack: 4
seed: 0
total_timesteps: 7000000      # match P1 (extended from 2M)
eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 1.0e-4
  buffer_size: 500000          # match P1 (~28GB; sized for shared server RAM)
  learning_starts: 100000
  batch_size: 32
  gamma: 0.99
  train_freq: 4
  gradient_steps: 1
  target_update_interval: 1000
  exploration_fraction: 0.5    # <-- slow decay (the ablation)

```

```
exploration_initial_eps: 1.0
exploration_final_eps: 0.01
```

7.3.5 configs/P4_buffersmall.yaml (30 lines)

```
# P4 - DQN with a SMALL replay buffer (Tier B).
# Aspect: memory / sample correlation. A small buffer keeps only very recent,
# highly correlated transitions, weakening the i.i.d. assumption behind
# experience replay. Compare against P1 (buffer_size = 200000).
exp_id: P4_buffersmall
task: pong
algo: dqn
env_id: ALE/Pong-v5
policy: CnnPolicy
n_envs: 1
n_stack: 4
seed: 0
total_timesteps: 2000000      # T2 ablation budget; compare vs P1's 2M snapshot from metri
eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 1.0e-4
  buffer_size: 20000          # <-- small buffer (the ablation)
  learning_starts: 10000     # must be <= buffer_size
  batch_size: 32
  gamma: 0.99
  train_freq: 4
  gradient_steps: 1
  target_update_interval: 1000
  exploration_fraction: 0.1
  exploration_initial_eps: 1.0
  exploration_final_eps: 0.01
```

7.3.6 configs/P5_ppo_pong.yaml (29 lines)

```
# P5 - PPO on Pong (Tier B): policy-based arm of the algorithm-family thread.
# Aspect: algorithm family (value-based DQN vs policy-based PPO on the SAME task).
# Pair with P1 in the report. PPO needs many more frames than DQN on Atari, so
# total_timesteps is large; reduce it if compute/time is short.
exp_id: P5_ppo_pong
task: pong
algo: ppo
```



```

env_id: ALE/Pong-v5
policy: CnnPolicy
n_envs: 16 # SubprocVecEnv parallel rollouts (~16 CPU threads)
n_stack: 4
seed: 0
total_timesteps: 7000000 # equalized with P1 (DQN) for a fair T1 comparison
# (same env-step budget; Pong converges < 7M for PPO)

eval_freq: 250000
eval_episodes: 20
checkpoint_freq: 1000000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 2.5e-4
  n_steps: 128 # rollout length per env (16 * 128 = 2048 / update)
  batch_size: 512
  n_epochs: 4
  gamma: 0.99
  gae_lambda: 0.95
  clip_range: 0.1
  ent_coef: 0.01
  vf_coef: 0.5

```

7.3.7 configs/P5b_ppo_zoo.yaml (37 lines)

```

# P5b - PPO on Pong with the SB3-zoo / Schulman-2017 Atari recipe.
# Sanity check for T1: our original P5 used constant LR / clip_range and lost
# (-6.58 at 7M, far below literature). Literature PPO on Pong reaches +18-21
# at ~7-10M with LINEAR decay on both learning_rate and clip_range. This run
# applies the literature recipe to isolate whether T1's "DQN beats PPO" surface
# result was a hyperparameter artifact.
#
# Differences vs P5 (the originally-published numbers):
#   learning_rate: 2.5e-4 -> lin_2.5e-4 (linear decay to 0)
#   clip_range:    0.1    -> lin_0.1    (linear decay to 0)
#   batch_size:    512    -> 256       (zoo default; smaller minibatches)
# Everything else (n_envs, n_steps, n_epochs, gamma, ent_coef, vf_coef) matches
# P5 so the only thing being isolated is the LR / clip / batch-size recipe.
exp_id: P5b_ppo_zoo
task: pong
algo: ppo
env_id: ALE/Pong-v5
policy: CnnPolicy
n_envs: 16
n_stack: 4

```

```

seed: 0
total_timesteps: 7000000      # same budget as P5 (fair T1-side comparison)
eval_freq: 250000
eval_episodes: 20
checkpoint_freq: 1000000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: lin_2.5e-4    # SB3 linear schedule: 2.5e-4 -> 0
  n_steps: 128
  batch_size: 256             # zoo default
  n_epochs: 4
  gamma: 0.99
  gae_lambda: 0.95
  clip_range: lin_0.1         # SB3 linear schedule: 0.1 -> 0
  ent_coef: 0.01
  vf_coef: 0.5

```

7.3.8 configs/V1_defendcenter.yaml (32 lines)

```

# V1 - VizDoom Defend-the-Center, PPO (Tier A): the 3D headline baseline.
# Reference for Task 2. Also the reused arm for:
#   V2 = "Discrete actions", V4 = "frame stack 4", V5 = "PPO".
# Run >=3 seeds: --seed 0 / 1 / 2
# NOTE: confirm env id on first run (VizdoomDefendCenter-v0 expected).
exp_id: V1_defendcenter
task: vizdoom
algo: ppo
env_id: VizdoomDefendCenter-v1
policy: CnnPolicy
n_envs: 16
n_stack: 4
frame_skip: 4
obs_shape: [84, 84]
grayscale: true
seed: 0
total_timesteps: 2000000      # ~30-90 min on an RTX 4090
eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 2.5e-4
  n_steps: 128

```

```

batch_size: 512
n_epochs: 4
gamma: 0.99
gae_lambda: 0.95
clip_range: 0.2
ent_coef: 0.01
vf_coef: 0.5

```

7.3.9 configs/V2_multibinary.yaml (35 lines)

```

# V2 - VizDoom Defend-the-Center, PPO, MULTI-button action space (Tier A).
# Aspect: action-space design (unique to VizDoom). The Farama wrapper exposes a
# Discrete space of single-button presses by default (max_buttons_pressed = 1).
# Setting max_buttons_pressed = 0 lets the agent press any combination of
# buttons at once -> a MultiDiscrete ("multi-binary"-style) action space.
# Compare against V1 (Discrete). PPO supports MultiDiscrete; DQN would not.
exp_id: V2_multibinary
task: vizdoom
algo: ppo
env_id: VizdoomDefendCenter-v1
policy: CnnPolicy
n_envs: 16
n_stack: 4
frame_skip: 4
obs_shape: [84, 84]
grayscale: true
env_kwargs:
  max_buttons_pressed: 0      # <-- multi-button (MultiDiscrete) action space
seed: 0
total_timesteps: 2000000
eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 2.5e-4
  n_steps: 128
  batch_size: 512
  n_epochs: 4
  gamma: 0.99
  gae_lambda: 0.95
  clip_range: 0.2
  ent_coef: 0.01
  vf_coef: 0.5

```

7.3.10 configs/V3_healthgathering.yaml (31 lines)

```
# V3 - VizDoom Health Gathering, PPO (Tier B): hardest scenario in the ladder.
# Aspect: task complexity. Used in the difficulty-scaling thread together with
# V0 (Basic) and V1 (Defend-Center): final reward vs scenario difficulty.
# NOTE: confirm env id on first run (VizdoomHealthGathering-v0 expected).
exp_id: V3_healthgathering
task: vizdoom
algo: ppo
env_id: VizdoomHealthGathering-v1
policy: CnnPolicy
n_envs: 16
n_stack: 4
frame_skip: 4
obs_shape: [84, 84]
grayscale: true
seed: 0
total_timesteps: 3000000      # harder exploration -> larger budget
eval_freq: 150000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 2.5e-4
  n_steps: 128
  batch_size: 512
  n_epochs: 4
  gamma: 0.99
  gae_lambda: 0.95
  clip_range: 0.2
  ent_coef: 0.01
  vf_coef: 0.5
```

7.3.11 configs/V4_stack1.yaml (31 lines)

```
# V4 - VizDoom Defend-the-Center, PPO, NO frame stacking (Tier B).
# Aspect: partial observability. A single frame (n_stack = 1) hides motion and
# direction, which matters more in a 3D first-person view than in 2D Pong.
# Compare against V1 (n_stack = 4).
exp_id: V4_stack1
task: vizdoom
algo: ppo
env_id: VizdoomDefendCenter-v1
policy: CnnPolicy
```

```

n_envs: 16
n_stack: 1                                # <-- no temporal context (the ablation)
frame_skip: 4
obs_shape: [84, 84]
grayscale: true
seed: 0
total_timesteps: 2000000
eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true
algo_kwargs:
  learning_rate: 2.5e-4
  n_steps: 128
  batch_size: 512
  n_epochs: 4
  gamma: 0.99
  gae_lambda: 0.95
  clip_range: 0.2
  ent_coef: 0.01
  vf_coef: 0.5

```

7.3.12 configs/V5_dqn_defendcenter.yaml (32 lines)

```

# V5 - VizDoom Defend-the-Center, DQN (Tier B): value-based arm, mirror of P5.
# Aspect: algorithm family (DQN vs PPO on the SAME 3D task). Pair with V1 in the
# report. Keeps the default Discrete action space (DQN requires it).
exp_id: V5_dqn_defendcenter
task: vizdoom
algo: dqn
env_id: VizdoomDefendCenter-v1
policy: CnnPolicy
n_envs: 1
n_stack: 4
frame_skip: 4
obs_shape: [84, 84]
grayscale: true
seed: 0
total_timesteps: 2000000
eval_freq: 100000
eval_episodes: 20
checkpoint_freq: 500000
device: cuda
record_gif: true

```

```
algo_kwargs:
  learning_rate: 1.0e-4
  buffer_size: 500000          # Large buffer (~28GB; sized for shared server RAM)
  learning_starts: 50000
  batch_size: 32
  gamma: 0.99
  train_freq: 4
  gradient_steps: 1
  target_update_interval: 1000
  exploration_fraction: 0.2
  exploration_initial_eps: 1.0
  exploration_final_eps: 0.05
```

7.4 B.4 Analysis pipeline (figures, summary tables, this appendix)

7.4.1 tools/build_report_assets.py (285 lines)

"""Build the report's grouped T1-T5 figures + numbered summary tables.

Reads ``results/\*/eval.json`` and ``results/\*/metrics.csv`` after a fetch, and writes everything the report needs into ``report\_assets/``:

T1_algo_family.png	Pong DQN vs PPO + Defend-Center PPO vs DQN (subplots)
T2_dqn_components.png	Pong DQN ablations (baseline, target-off, eps, buffer)
T3_framestack.png	VizDoom Defend-Center stack-4 vs stack-1
T4_actionspace.png	VizDoom Defend-Center Discrete vs MultiDiscrete
T5_difficulty.png	Final reward by VizDoom scenario (DC, HG)
summary.md	Auto-filled tables for RESULTS.md / the report

Experiments without any results are skipped gracefully (the script just prints which ones it couldn't build). Run after ``pwsh deploy/fetch.ps1 -Lite``:

```
python tools/build_report_assets.py
"""
from __future__ import annotations

import glob
import json
import statistics
import sys
from pathlib import Path

# This repo's path contains CJK characters; force UTF-8 stdout so our prints
# don't crash on the Windows cp1252 console.
try:
    sys.stdout.reconfigure(encoding="utf-8", errors="replace")
    sys.stderr.reconfigure(encoding="utf-8", errors="replace")
except Exception: # noqa: BLE001
    pass

import matplotlib

matplotlib.use("Agg") # headless-safe
import matplotlib.pyplot as plt # noqa: E402
import numpy as np # noqa: E402

# Make ``common`` importable when run as a script.
sys.path.insert(0, str(Path(__file__).resolve().parents[1]))
```

```

from common.plotting import aggregate_seeds # noqa: E402
from common.utils import REPO_ROOT, RESULTS_DIR # noqa: E402

ASSETS = REPO_ROOT / "report_assets"
ASSETS.mkdir(parents=True, exist_ok=True)

# -----
# Aggregation helpers
# -----

def eval_summary(exp_id: str):
    """(n_seeds, mean, std-across-seed-means, [per-seed mean rewards])."""
    rewards = []
    for f in sorted(glob.glob(str(RESULTS_DIR / f"{exp_id}_s*" / "eval.json"))):
        with open(f, encoding="utf-8") as fh:
            rewards.append(json.load(fh)["mean_reward"])
    if not rewards:
        return 0, float("nan"), float("nan"), []
    mean = statistics.mean(rewards)
    std = statistics.pstdev(rewards) if len(rewards) > 1 else 0.0
    return len(rewards), mean, std, rewards

def fmt(stats):
    n, m, s, _ = stats
    return "—" if n == 0 else f"{m:.2f} ± {s:.2f} (n={n})"

def overlay_curves(exp_ids, labels, ax, title):
    """Plot mean±std learning curves for several experiments on one axis."""
    plotted_any = False
    for exp_id, label in zip(exp_ids, labels):
        try:
            x, mean, std = aggregate_seeds(exp_id)
            ax.plot(x, mean, label=label)
            ax.fill_between(x, mean - std, mean + std, alpha=0.2)
            plotted_any = True
        except Exception as exc: # noqa: BLE001
            print(f" skip {exp_id}: {exc}")
    ax.set_xlabel("Timesteps")
    ax.set_ylabel("Episode reward (mean ± std across seeds)")
    ax.set_title(title)
    ax.grid(alpha=0.3)
    if plotted_any:

```



```

        ax.legend()
    return plotted_any

# -----
# Per-thread figure builders
# -----
def build_T1():
    """T1a now overlays THREE Pong curves: DQN baseline, original PPO, and the
    SB3-zoo / Schulman 2017 literature-recipe PPO (linear LR + clip decay).
    The zoo recipe is the sanity-check for whether the DQN > PPO inversion at
    7 M is a hyperparameter artifact (it isn't — both PPO variants lose to DQN).
    """
    fig, axes = plt.subplots(1, 2, figsize=(11.5, 4.5))
    overlay_curves(
        ["P1", "P5_ppo_pong", "P5b_ppo_zoo"],
        ["DQN (P1)", "PPO original (P5)", "PPO zoo recipe (P5b)"],
        axes[0],
        "T1a — Pong: DQN vs PPO (two hyperparam recipes)",
    )
    overlay_curves(
        ["V1_defendcenter", "V5_dqn_defendcenter"],
        ["PPO (V1)", "DQN (V5)"],
        axes[1],
        "T1b — Defend-Center: policy-based vs value-based",
    )
    fig.tight_layout()
    out = ASSETS / "T1_algo_family.png"
    fig.savefig(out, dpi=150)
    plt.close(fig)
    return out

def build_T2():
    fig, ax = plt.subplots(figsize=(8.5, 5))
    overlay_curves(
        ["P1", "P2_targetoff", "P3_epsfast", "P3_epsslow", "P4_buffersmall"],
        ["Baseline (P1)", "Target net OFF (P2)", "ε fast decay (P3a)",
         "ε slow decay (P3b)", "Small buffer (P4)"],
        ax,
        "T2 — DQN components on Pong (2 M ablation budget)",
    )
    # Clip to the 2 M ablation budget. P1 was additionally extended to 7 M for
    # T1, and one P2 seed completed 7 M, but T2 is anchored at 2 M for a clean
    # 3-seed comparison across all ablations.

```

```

ax.set_xlim(0, 2.05e6)
fig.tight_layout()
out = ASSETS / "T2_dqn_components.png"
fig.savefig(out, dpi=150)
plt.close(fig)
return out

```

```

def build_T3():
    fig, ax = plt.subplots(figsize=(7.5, 4.5))
    overlay_curves(
        ["V1_defendcenter", "V4_stack1"],
        ["Stack 4 (V1)", "Stack 1 (V4)"],
        ax,
        "T3 — Frame stacking on VizDoom Defend-Center",
    )
    fig.tight_layout()
    out = ASSETS / "T3_framestack.png"
    fig.savefig(out, dpi=150)
    plt.close(fig)
    return out

```

```

def build_T4():
    fig, ax = plt.subplots(figsize=(7.5, 4.5))
    overlay_curves(
        ["V1_defendcenter", "V2_multibinary"],
        ["Discrete (V1)", "MultiDiscrete (V2)"],
        ax,
        "T4 — Action-space design on VizDoom Defend-Center",
    )
    fig.tight_layout()
    out = ASSETS / "T4_actionspace.png"
    fig.savefig(out, dpi=150)
    plt.close(fig)
    return out

```

```

def build_T5():
    """Final-reward bar chart across difficulty (DC, HG). Different scenarios
    have different reward scales, so a bar chart of final scores is clearer
    than overlaying learning curves on a shared y-axis. Basic (V0) is dropped.
    """
    exps = ["V1_defendcenter", "V3_healthgathering"]
    labels = ["Defend-Center (V1)", "Health-Gathering (V3)"]

```

```

means, stds, ns = [], [], []
for e in exps:
    n, m, s, _ = eval_summary(e)
    means.append(m if n else 0)
    stds.append(s if n else 0)
    ns.append(n)
fig, ax = plt.subplots(figsize=(6.5, 4.5))
bars = ax.bar(labels, means, yerr=stds, capsize=6,
              color=["tab:blue", "tab:orange"])
ax.set_ylabel("Final eval reward (mean of seed means)")
ax.set_title("T5 — Difficulty ladder (final reward by scenario)")
ax.grid(alpha=0.3, axis="y")
for bar, m, n in zip(bars, means, ns):
    label = f"{m:.1f} (n={n})" if n else "no data"
    ax.text(bar.get_x() + bar.get_width() / 2,
            m + (0.02 * max(abs(np.array(means)).max(), 1)),
            label, ha="center", va="bottom")
fig.tight_layout()
out = ASSETS / "T5_difficulty.png"
fig.savefig(out, dpi=150)
plt.close(fig)
return out

```

```

# -----
# Markdown summary
# -----

```

```

def build_summary_md():
    p1 = eval_summary("P1")
    p2 = eval_summary("P2_targetoff")
    p3a = eval_summary("P3_epsfast")
    p3b = eval_summary("P3_epsslow")
    p4 = eval_summary("P4_buffersmall")
    p5 = eval_summary("P5_ppo_pong")
    p5b = eval_summary("P5b_ppo_zoo")
    v1 = eval_summary("V1_defendcenter")
    v2 = eval_summary("V2_multibinary")
    v3 = eval_summary("V3_healthgathering")
    v4 = eval_summary("V4_stack1")
    v5 = eval_summary("V5_dqn_defendcenter")

    L = []
    L.append("# Results summary (auto-generated)\n")
    L.append("Numbers are `mean ± std` of per-seed eval means. `n` = seeds used.\n")

```

```

L.append("\n## T1 — Algorithm family (DQN vs PPO, equal env-step budget)\n")
L.append("| Task | DQN | PPO |")
L.append("|---|---|---|")
L.append(f"| Pong (7M, original PPO recipe) | {fmt(p1)} | {fmt(p5)} |")
L.append(f"| Pong (7M, SB3-zoo PPO recipe) | {fmt(p1)} | {fmt(p5b)} |")
L.append(f"| Defend-Center (2M each) | {fmt(v5)} | {fmt(v1)} |")

L.append("\n## T2 — DQN components on Pong\n")
L.append("| Variant | Reward |")
L.append("|---|---|")
L.append(f"| Baseline (P1) | {fmt(p1)} |")
L.append(f"| Target net OFF (P2) | {fmt(p2)} |")
L.append(f"|  $\epsilon$  fast (P3a) | {fmt(p3a)} |")
L.append(f"|  $\epsilon$  slow (P3b) | {fmt(p3b)} |")
L.append(f"| Small buffer (P4) | {fmt(p4)} |")

L.append("\n## T3 — Partial observability (frame stacking, Defend-Center)\n")
L.append("| Stack | Reward |")
L.append("|---|---|")
L.append(f"| 4 (V1) | {fmt(v1)} |")
L.append(f"| 1 (V4) | {fmt(v4)} |")

L.append("\n## T4 — Action-space design (Defend-Center)\n")
L.append("| Action space | Reward |")
L.append("|---|---|")
L.append(f"| Discrete (V1) | {fmt(v1)} |")
L.append(f"| MultiDiscrete (V2) | {fmt(v2)} |")

L.append("\n## T5 — Difficulty ladder (Basic dropped)\n")
L.append("| Scenario | Reward |")
L.append("|---|---|")
L.append(f"| Defend-Center (V1) | {fmt(v1)} |")
L.append(f"| Health-Gathering (V3) | {fmt(v3)} |")

out = ASSETS / "summary.md"
out.write_text("\n".join(L) + "\n", encoding="utf-8")
return out

def main():
    builders = [
        ("T1 algo family", build_T1),
        ("T2 DQN components", build_T2),
        ("T3 frame stack", build_T3),
        ("T4 action space", build_T4),

```

```

    ("T5 difficulty", build_T5),
    ("summary.md", build_summary_md),
]
for name, fn in builders:
    try:
        out = fn()
        # Print a repo-relative path so the CJK-laden absolute path doesn't
        # leak into the console output.
        try:
            rel = Path(out).relative_to(REPO_ROOT)
        except ValueError:
            rel = Path(out).name
        print(f"[ok]    {name:20s} -> {rel}")
    except Exception as exc: # noqa: BLE001
        print(f"[skip]  {name:20s} ({exc})")

if __name__ == "__main__":
    main()

```

7.4.2 tools/build_appendix.py (164 lines)

"""Auto-generate report/APPENDIX_CODE.md from the source tree.

*The course requirement is: *include your program code as an appendix (not counting toward the 10-page limit)*. Rather than hand-paste files into the report -- which goes stale the moment any source file is edited -- this script walks a curated list of paths and emits one Markdown file with the content of each, wrapped in fenced code blocks with the right language tag.*

Usage::

```

python tools/build_appendix.py [output_path]
# defaults to report/APPENDIX_CODE.md if no path is given

```

That generated file is then appended to ``report/REPORT.md`` (or included at pandoc time) so the PDF carries the full, current source code as Appendix B. We write directly to a UTF-8 file (not stdout) because the Windows console's cp1252 codec cannot encode characters like the ``ε`` literal in our code.

Grouping is by role so a reader can navigate top-down:

1. Common engine (the algorithm-agnostic training/eval stack)
2. Per-task entry (train_*.py and eval_record.py for each task/algo)
3. Experiment configs (one YAML per experiment - the experiment design)

- 4. Analysis pipeline (tools/* -- figures, summary tables, this appendix)
- 5. Deployment helpers (selected SSH/scp scripts for the no-git server)

Files NOT included on purpose:

- common/__init__.py empty package marker
- smoke_test.py dev-time only, not part of experiments
- deploy/*.ps1 Windows-specific workflow plumbing

```
from __future__ import annotations
```

```
import sys
```

```
from pathlib import Path
```

```
REPO = Path(__file__).resolve().parents[1]
```

```
def lang_for(path: Path) -> str:
```

```
    """Best-effort fenced-code-block language tag for the file's extension."""
```

```
    suf = path.suffix.lower()
```

```
    return {
```

```
        ".py": "python",
```

```
        ".yaml": "yaml",
```

```
        ".yml": "yaml",
```

```
        ".sh": "bash",
```

```
        ".ps1": "powershell",
```

```
        ".md": "markdown",
```

```
        ".json": "json",
```

```
    }.get(suf, "text")
```

```
# Order matters: this is the order the appendix presents the project in.
```

```
GROUPS = [
```

```
    (
```

```
        "B.1 Common engine (algorithm-agnostic training & evaluation stack)",
```

```
        [
```

```
            "common/utils.py",
```

```
            "common/envs.py",
```

```
            "common/vizdoom_wrappers.py",
```

```
            "common/callbacks.py",
```

```
            "common/eval_utils.py",
```

```
            "common/plotting.py",
```

```
            "common/train_core.py",
```

```
        ],
```

```
    ),
```

```
    (
```

```

    "B.2  Per-task training & evaluation entry points",
    [
        "pong/train_pong_dqn.py",
        "pong/train_pong_ppo.py",
        "vizdoom/train_vizdoom_dqn.py",
        "vizdoom/train_vizdoom_ppo.py",
        "pong/eval_record.py",
        "vizdoom/eval_record.py",
    ],
),
(
    "B.3  Experiment configs (one YAML per (task, algorithm, ablation))",
    [
        "configs/P1.yaml",
        "configs/P2_targetoff.yaml",
        "configs/P3_epsfast.yaml",
        "configs/P3_epsslow.yaml",
        "configs/P4_buffersmall.yaml",
        "configs/P5_ppo_pong.yaml",
        "configs/P5b_ppo_zoo.yaml",
        "configs/V1_defendcenter.yaml",
        "configs/V2_multibinary.yaml",
        "configs/V3_healthgathering.yaml",
        "configs/V4_stack1.yaml",
        "configs/V5_dqn_defendcenter.yaml",
    ],
),
(
    "B.4  Analysis pipeline (figures, summary tables, this appendix)",
    [
        "tools/build_report_assets.py",
        "tools/build_appendix.py",
    ],
),
(
    "B.5  Deployment helpers (selected SSH/scp scripts)",
    [
        "deploy/_activate.sh",
        "deploy/remote_setup.sh",
        "deploy/remote_train.sh",
        "deploy/remote_queue.sh",
        "deploy/remote_queue_round3.sh",
    ],
),
]

```

```
def emit_header(out):
    out.write(
        """# Appendix B — Program code
```

The full source tree is reproduced below, grouped by role for top-down reading. The same files are also available at <<https://github.com/Ashurali/RLAIHW3>> (branch `main`); this appendix is a self-contained, frozen snapshot for grading.

Every source file is auto-extracted from the repository by ``tools/build_appendix.py``, so what you read here is exactly what trained the models behind the results in §3. The code is documented inline — each module opens with a docstring explaining its role, and non-obvious decisions (schedule resolver, frame-stack wrapper, CSV-corruption fallback, GPU perf flags) carry comments at point-of-use.

```
"""
    )
```

```
def emit_file(out, rel_path: str):
    path = REPO / rel_path
    if not path.exists():
        out.write(f"### `{rel_path}`\n\n*(file missing at appendix-build time — skipped)*")
        return
    content = path.read_text(encoding="utf-8")
    # Strip trailing whitespace on lines; keep blank lines.
    content = "\n".join(line.rstrip() for line in content.splitlines())
    n_lines = len(content.splitlines())
    lang = lang_for(path)
    out.write(f"### `{rel_path}` ({n_lines} lines)\n\n")
    out.write(f"```{lang}\n{content}\n```\n\n")
```

```
def main():
    out_path = Path(sys.argv[1]) if len(sys.argv) > 1 else REPO / "report" / "APPENDIX_CODE"
    out_path.parent.mkdir(parents=True, exist_ok=True)
    with out_path.open("w", encoding="utf-8", newline="\n") as out:
        emit_header(out)
        for title, files in GROUPS:
            out.write(f"\n\nnewpage\n\n## {title}\n\n")
            for rel in files:
                emit_file(out, rel)
    # Lightweight summary to stdout (ASCII only, safe for Windows console).
```



```
n_lines = sum(1 for _ in out_path.open(encoding="utf-8"))
print(f"wrote {out_path.relative_to(REPO)} ({n_lines} lines)")

if __name__ == "__main__":
    main()
```

7.5 B.5 Deployment helpers (selected SSH/scp scripts)

7.5.1 deploy/_activate.sh (68 lines)

```
#!/usr/bin/env bash
# Sourced by the remote deploy scripts to create/activate the training env.
# Supports two backends so it works with or without sudo:
#   ENV_KIND=venv    -> python -m venv (needs python3.x-venv / ensurepip)
#   ENV_KIND=conda   -> conda env      (no system packages / sudo needed)
#
# Reads (with defaults) from the environment, passed by deploy/remote.ps1:
#   ENV_KIND          conda | venv          (default: conda)
#   VENV              conda env name OR venv dir (default: hw3 / REMOTE_VENV)
#   PY_VERSION        python for the conda env (default: 3.12)
#   PY_BIN             interpreter for venv creation (auto-detect if empty)
#   REMOTE_CONDA_BASE conda base dir override   (auto-detect if empty)

VENV="${VENV:-${REMOTE_VENV:-hw3}}"
ENV_KIND="${ENV_KIND:-conda}"
PY_VERSION="${PY_VERSION:-3.12}"
PY_BIN="${PY_BIN:-}"

_ensure_conda() {
    command -v conda >/dev/null 2>&1 && return 0
    local c
    for c in "${REMOTE_CONDA_BASE:-}" "$HOME/miniconda3" "$HOME/anaconda3" \
        "$HOME/miniforge3" "$HOME/mambaforge" /opt/conda /opt/anaconda3; do
        if [ -n "$c" ] && [ -f "$c/etc/profile.d/conda.sh" ]; then
            # shellcheck disable=SC1091
            source "$c/etc/profile.d/conda.sh"
            return 0
        fi
    done
    echo "ERROR: cannot locate conda. Set REMOTE_CONDA_BASE in deploy/server.env" >&2
    echo "        to the output of 'conda info --base' on the server." >&2
    return 1
}

_pick_py() {
    local p
    for p in python3.12 python3.11 python3.10 python3; do
        command -v "$p" >/dev/null 2>&1 && { echo "$p"; return 0; }
    done
    return 1
}
```

```

# Activate an already-created environment into the current shell.
activate_env() {
    if [ "$ENV_KIND" = "conda" ]; then
        _ensure_conda || return 1
        conda activate "$VENV"
    else
        # shellcheck disable=SC1091
        conda activate "$VENV"
    fi
}

# Create the environment, then activate it.
create_env() {
    if [ "$ENV_KIND" = "conda" ]; then
        _ensure_conda || return 1
        echo "conda env: $VENV (python=$PY_VERSION) under $(conda info --base)"
        conda create -y -n "$VENV" "python=$PY_VERSION"
        conda activate "$VENV"
    else
        local py="${PY_BIN:-$_pick_py}"
        echo "venv interpreter: $py = $($py --version 2>&1)"
        "$py" -m venv "$VENV"
        # shellcheck disable=SC1091
        conda activate $VENV
    fi
}

```

7.5.2 deploy/remote_setup.sh (55 lines)

```

#!/usr/bin/env bash
# One-time server setup, DETACHED (survives SSH Logout). Probes the hardware,
# creates the conda env, installs requirements, and verifies CUDA.
#
# Usage (via deploy/remote.ps1 -Action setup, or directly on the server):
#   REMOTE_VENV=hw3 bash deploy/remote_setup.sh
set -euo pipefail
cd "$(dirname "$0")/.." # repo root on the server
VENV="${REMOTE_VENV:-hw3}"
ENV_KIND="${ENV_KIND:-conda}"
export REMOTE_VENV ENV_KIND PY_VERSION PY_BIN REMOTE_CONDA_BASE
mkdir -p logs

run_setup() {
    echo "=== server probe ($(date '+%F %T')) ==="

```

```

if command -v nvidia-smi >/dev/null 2>&1; then nvidia-smi; else echo "nvidia-smi not found"
echo "CPUs(threads): $(nproc 2>/dev/null || echo '?')"
echo "RAM: $(free -h 2>/dev/null | awk '/Mem:/{print $2}' || echo '?')"
echo "Python(default): $(python3 --version 2>&1)"
echo "Available interpreters:"
for p in python3.12 python3.11 python3.10; do
    command -v "$p" >/dev/null 2>&1 && echo "  - $p = $($p --version 2>&1)"
done

# Create the env via deploy/_activate.sh (conda or venv per ENV_KIND). conda
# needs no sudo / system packages, which is why it's the default on the lab box.
echo "=== creating environment ($ENV_KIND -> $VENV) ==="
source deploy/_activate.sh
set +u; create_env; set -u
python -m pip install --upgrade pip wheel

echo "=== installing requirements (CUDA torch via extra-index-url) ==="
pip install -r requirements.txt

echo "=== verifying torch sees the GPU ==="
python - <<'PY'
import torch
print("torch", torch.__version__, "| cuda_available:", torch.cuda.is_available())
if torch.cuda.is_available():
    print("device:", torch.cuda.get_device_name(0), "| count:", torch.cuda.device_count())
    print("cuda runtime:", torch.version.cuda)
else:
    print("WARNING: CUDA not available -- training will fall back to CPU.")
PY
echo "SETUP_DONE"
}

# Re-exec a detached worker so the install survives SSH Logout.
if [ "${1:-}" = "--worker" ]; then run_setup; exit 0; fi

log="logs/setup.log"
nohup bash "$0" --worker > "$log" 2>&1 < /dev/null &
echo $! > logs/setup.pid
echo "STARTED setup pid=$(cat logs/setup.pid)"
echo "  log: $log    (wait for SETUP_DONE: tail -f $log)"

```

7.5.3 deploy/remote_train.sh (36 lines)

```

#!/usr/bin/env bash
# Launch ONE training run DETACHED on the server (survives SSH Logout).

```

```
# Output goes to logs/<exp>_s<seed>.log so progress is trackable with tail -f.
#
# Usage (run from the repo root on the server, or via deploy/remote.ps1):
#   REMOTE_VENV=hw3 bash deploy/remote_train.sh configs/P1.yaml 0
set -euo pipefail
cd "$(dirname "$0")/.." # repo root on the server

CONFIG="${1:?usage: remote_train.sh <config.yaml> [seed]}"
SEED="${2:-0}"
VENV="${REMOTE_VENV:-hw3}"
mkdir -p logs

name="$(basename "$CONFIG" .yaml)"
task="$(awk '/^task:/{print $2; exit}' "$CONFIG")"
algo="$(awk '/^algo:/{print $2; exit}' "$CONFIG")"
case "${task}_${algo}" in
    pong_dqn)    script="pong/train_pong_dqn.py" ;;
    pong_ppo)    script="pong/train_pong_ppo.py" ;;
    vizdoom_ppo) script="vizdoom/train_vizdoom_ppo.py" ;;
    vizdoom_dqn) script="vizdoom/train_vizdoom_dqn.py" ;;
    *) echo "ERROR: cannot map task='$task' algo='$algo' from $CONFIG" && exit 1 ;;
esac

log="logs/${name}_s${SEED}.log"
pidf="logs/${name}_s${SEED}.pid"

# nohup + redirect + </dev/null detaches the job so it keeps running after the
# SSH session closes. -u keeps Python output unbuffered for live tailing.
nohup bash -c "source deploy/_activate.sh && activate_env && exec python -u '$script' --co
    > "$log" 2>&1 < /dev/null &
echo $! > "$pidf"

echo "STARTED $name seed=$SEED pid=$(cat "$pidf") script=$script"
echo "  log: $log    (track: tail -f $log)"
```

7.5.4 deploy/remote_queue.sh (67 lines)

```
#!/usr/bin/env bash
# Run a LIST of experiments SEQUENTIALLY, detached (survives SSH logout).
# One GPU -> run jobs back-to-back rather than thrashing in parallel.
# Master Log: logs/queue_<timestamp>.log ; per-run: logs/<exp>_s<seed>.log
#
# Usage (via deploy/remote.ps1 -Action queue, or directly):
#   REMOTE_VENV=hw3 bash deploy/remote_queue.sh
set -euo pipefail
```

```

cd "$(dirname "$0")/.."                                # repo root on the server
VENV="${REMOTE_VENV:-hw3}"
mkdir -p logs

# --- EDIT THIS QUEUE -----
# Each entry is "<config_basename> <seed>". Tier A x3 seeds shown by default;
# uncomment Tier B lines to queue the full scope.
LIST=(
  # Round 3: catch-up run on 2026-05-29 to fill the T1b gap.
  # ONLY V5 (DQN on Defend-Center), 3 seeds, 2M steps. Everything else is done.
  "V5_dqn_defendcenter 0" "V5_dqn_defendcenter 1" "V5_dqn_defendcenter 2"
)
# -----

map_script() {
  local cfg="$1" task algo
  task="$(awk '/^task:/{print $2; exit}' "$cfg")"
  algo="$(awk '/^algo:/{print $2; exit}' "$cfg")"
  case "${task}_${algo}" in
    pong_dqn)    echo "pong/train_pong_dqn.py" ;;
    pong_ppo)    echo "pong/train_pong_ppo.py" ;;
    vizdoom_ppo) echo "vizdoom/train_vizdoom_ppo.py" ;;
    vizdoom_dqn) echo "vizdoom/train_vizdoom_dqn.py" ;;
    *) return 1 ;;
  esac
}

run_queue() {
  source deploy/_activate.sh
  activate_env
  local n=${#LIST[@]} i=0
  for entry in "${LIST[@]"; do
    i=$((i + 1))
    # shellcheck disable=SC2086
    set -- $entry
    local name="$1" seed="$2" cfg="configs/$1.yaml" script rlog
    if [ ! -f "$cfg" ]; then echo "[${i}/${n}] SKIP $name: $cfg not found"; continue; fi
    if ! script="$(map_script "$cfg)"; then echo "[${i}/${n}] SKIP $name: cannot map task/alg";
    rlog="logs/${name}_s${seed}.log"
    echo "=== [${i}/${n}] $(date '+%F %T') START $name seed=$seed -> $rlog ==="
    if python -u "$script" --config "$cfg" --seed "$seed" > "$rlog" 2>&1; then
      echo "=== [${i}/${n}] $(date '+%F %T') DONE $name seed=$seed ==="
    else
      echo "=== [${i}/${n}] $(date '+%F %T') FAIL $name seed=$seed (see $rlog) ==="
    fi
  done
}

```

```

done
echo "QUEUE_DONE"
}

# Re-exec a detached worker so the whole queue survives SSH Logout.
if [ "${1:-}" = "--worker" ]; then run_queue; exit 0; fi

ts="$(date +%Y%m%d_%H%M%S)"
qlog="logs/queue_${ts}.log"
nohup bash "$0" --worker > "$qlog" 2>&1 < /dev/null &
echo $! > logs/queue.pid
echo "STARTED queue (${#LIST[@]} runs) pid=$(cat logs/queue.pid)"
echo "  master log: $qlog      (track:  tail -f $qlog)"
echo "  per-run logs: logs/<exp>_s<seed>.log"

```

7.5.5 deploy/remote_queue_round3.sh (63 lines)

```

#!/usr/bin/env bash
# Round 3 catch-up queue: runs AFTER V5 finishes.
#   - P4_buffersmall x 3 seeds at 2M (T2 small-buffer gap)
#   - P5b_ppo_zoo x 3 seeds at 7M (T1 PPO literature-recipe sanity check)
#
# Same nohup / detached pattern as remote_queue.sh. Launched by the V5 watcher
# when the round-2 queue master exits.
set -euo pipefail
cd "$(dirname "$0")/.." # repo root on the server
VENV="${REMOTE_VENV:-hw3}"
mkdir -p logs

LIST=(
  # Cheap small-buffer ablation first (2M, single-env DQN, ~30-60min/seed).
  "P4_buffersmall 0" "P4_buffersmall 1" "P4_buffersmall 2"
  # Literature-recipe PPO last (7M, 16 envs, ~2h/seed at idle GPU).
  "P5b_ppo_zoo 0" "P5b_ppo_zoo 1" "P5b_ppo_zoo 2"
)

map_script() {
  local cfg="$1" task algo
  task="$(awk '/^task:/{print $2; exit}' "$cfg")"
  algo="$(awk '/^algo:/{print $2; exit}' "$cfg")"
  case "${task}_${algo}" in
    pong_dqn)    echo "pong/train_pong_dqn.py" ;;
    pong_ppo)    echo "pong/train_pong_ppo.py" ;;
    vizdoom_ppo) echo "vizdoom/train_vizdoom_ppo.py" ;;
    vizdoom_dqn) echo "vizdoom/train_vizdoom_dqn.py" ;;
  esac
}

```

```

    *) return 1 ;;
esac
}

run_queue() {
    source deploy/_activate.sh
    activate_env
    local n=${#LIST[@]} i=0
    for entry in "${LIST[@]}; do
        i=$((i + 1))
        # shellcheck disable=SC2086
        set -- $entry
        local name="$1" seed="$2" cfg="configs/$1.yaml" script rlog
        if [ ! -f "$cfg" ]; then echo "[i/$n] SKIP $name: $cfg not found"; continue; fi
        if ! script="$(map_script "$cfg)"; then echo "[i/$n] SKIP $name: cannot map task/alg";
        rlog="logs/${name}_s${seed}.log"
        echo "=== [i/$n] $(date '+%F %T') START $name seed=$seed -> $rlog ==="
        if python -u "$script" --config "$cfg" --seed "$seed" > "$rlog" 2>&1; then
            echo "=== [i/$n] $(date '+%F %T') DONE $name seed=$seed ==="
        else
            echo "=== [i/$n] $(date '+%F %T') FAIL $name seed=$seed (see $rlog) ==="
        fi
    done
    echo "QUEUE_DONE"
}

if [ "${1:-}" = "--worker" ]; then run_queue; exit 0; fi

ts="$(date +%Y%m%d_%H%M%S)"
qlog="logs/queue_round3_${ts}.log"
nohup bash "$0" --worker > "$qlog" 2>&1 < /dev/null &
echo $! > logs/queue_round3.pid
echo "STARTED round3 queue (${#LIST[@]} runs) pid=$(cat logs/queue_round3.pid)"
echo "  master log: $qlog      (track: tail -f $qlog)"
echo "  per-run logs: logs/<exp>_s<seed>.log"

```